60001

Advanced Computer Architecture
Imperial College London

Contents

Chapter 1

Introduction

1.1 Course Structure and Logistics

Prof Paul Kelly

Teaching the entire course.

- Microprocessor design.
- Optimising software for hardware, and compiler design.
- Optimising hardware for specific software tasks.
- Challenges past, present & future.

Taught through pre-recorded lectures and live tutorial sessions.

This course is largely textbook based.

- 936 pages covering the course content and more.
- Useful appendices covering both introductory and advanced material.

The book is written by John Hennessy and David Patterson.

**Computer Architecture:
A Quantitative Approach (6th Edition)**

Chapter 1 - Part 1: Introduction

Chapter 2

Pipelining

MIPS/Microprocessor without Interlocked Pipelined Stages   **Definition 2.0.1**

MIPS is a reduced instruction set (RISC) architecture originally developed for the R2000 microprocessor.

- 3 types of instruction layouts
- Load-store architecture
- Support for floating point arithmetic

2.1 Instruction Layout

The instructions set architecture (ISA) determines the layout of instructions. Here we consider the mips architecture.

Register Type

31 26	25 21	20 16	15 11	10 6	5 0
opcode	Register Source 2 (Rs)	Register Source 1 (Rt)	Register Destination (Rd)	Shift code (shamt)	Function Code (funct)
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

5 bit register specifier
limits MIPS to 32 registers

Immediate Type

31 26	25 21	20 16	15 0
opcode	Register Source 2 (Rs)	Register Destination (Rd)	immediate operand
6 bits	5 bits	5 bits	16 bits

Opcode specifies how the fileds will be interpreted by specifying the instruction type   Opcode specifies how the fileds will be interpreted by specifying the instruction type

Jump Type

31 26	25 0
opcode	target
6 bits	26 bits

Can use word-alignment (4 bytes - 2 bits) of instructions to address with 28 bits

Branch (Immediate Type)

31 26	25 21	20 16	15 0
opcode	Register Source 2 (Rs)	Branch operation	pc relative address
6 bits	5 bits	5 bits	16 bits

address is offset from the current PC, with word alignment of instructions considered

The size of fields in the instruction layouts determines characteristics such as:

- Maximum number of registers
- Maximum distance for a conditional jump
- Size of immediate operands
- Range of addresses that can be used.

MIPS Assembly   *Extra Fun!* **2.1.1**

A basic guide listing of mips instructions can be found here Basic MIPS instructions.

2.2   Pipeline Structure

- Execution of an instruction is split into stages
- Throughput is potentially increased by factor $1/_{\text{number of stages}}$ (idealistically)
- All stages work on an instruction simultaneously/in parallel (very little extra hardware required for the speedup advantage)

The speedup is reduced by

- Latency increased due to latches
- Pipeline rate limited by slowest stage (unbalanced stages / fragmentation)
- Time required to fill and drain the pipeline.
- Pipeline hazards which result in stalls (unable to dispatch another instruction in a given cycle).

2.3   Pipeline Hazards

2.3.1   Structural Hazard

Structural Hazard	Definition 2.3.1
Where hardware is unable to support a combination of instructions.	

Multiple pipeline stages may need to access the same hardware resources:

- Register file (register operand fetch and register write back)
- Access to memory (RAM port in older machines, cache (SRAM) now)

Not enough ports!	Example Question 2.3.1

Given basic pipeline structure above, what structural hazard could occur between *Memory Access* and *instruction fetch* if there is only one RAM port?

- -

No instruction can be fetched when the *Memory Access* stage is filled, this results in a stall.

The maximum potential speedup for a 5 stage pipeline is $5\times$, however due to the stalls we can see a recurring pattern:

Cycle:	$6n$	$6n+1$	$6n+2$	$6n+3$	$6n+4$	$6n+5$
Instructions:	2	2	3	3	3	2

We would expect a $5\times$ speedup from this pipeline. However we are only getting a $2.5\times$ speedup due to the stalls.

2.3.2   Data Hazard

| **Data Hazard** | **Definition 2.3.2** |

Instruction is dependent on the result of a prior instruction still in the pipeline.

Most often caused by a dependency between instructions.

| **Forwarding Paths** | **Definition 2.3.3** |

Paths between pipeline stages to allow results from previous instructions (not yet written back) to be sent to instructions afterwards that are in the pipeline.

Result Used By Many Subsequent Instructions

```
li    $1, 100     # $1 = 100
addiu $2, $1, 100  # $1 += 100 # here onwards depends on $1
addiu $3, $1, 100  # $1 += 100
addiu $4, $1, 100  # $1 += 100
```
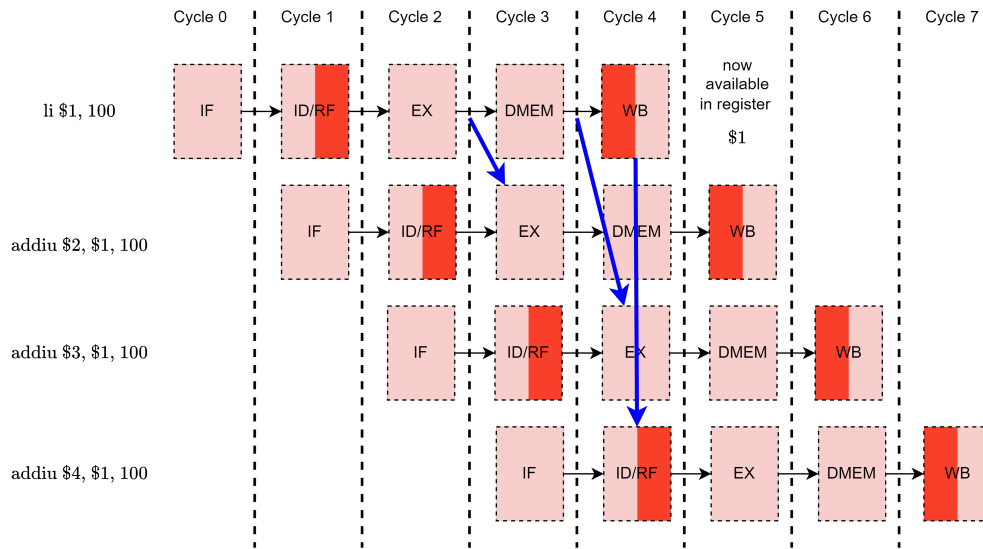
Chain of Results

```
li    $1, 100       # $1 = 100
addiu $2, $1, 100   # $1 += 100
addiu $3, $2, 100   # $1 += 100
addiu $4, $3, 100   # $1 += 100
```
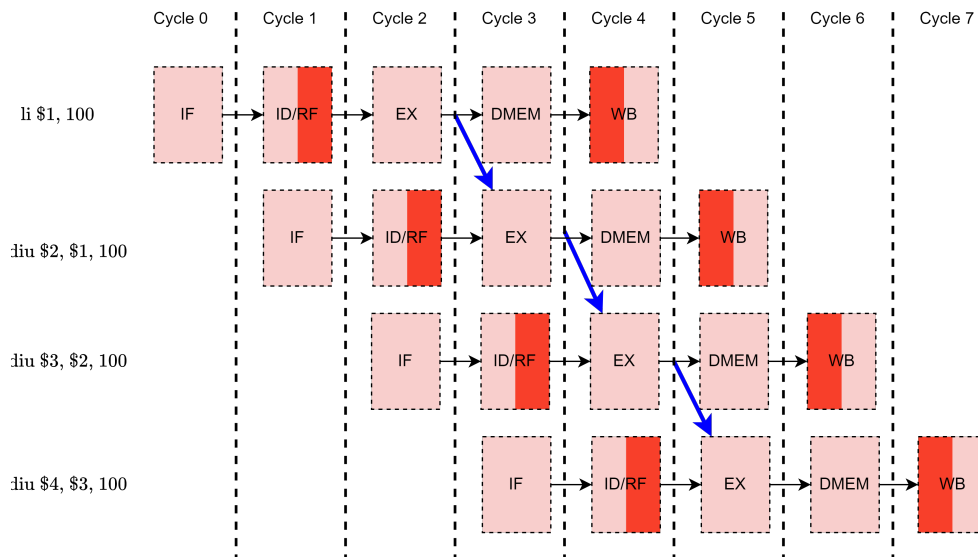

Data Hazard Despite Forwarding

Here have a *load to use stall/delay*, forwarding paths will not work here as the memory access stage is 2 stages later than execute (where the instruction is required).

```
lw    $1, 0($2)     # $1 = *($2)
addiu $2, $1, 100   # $1 += 100
addiu $3, $1, 100   # $1 += 100
addiu $4, $1, 100   # $1 += 100
```
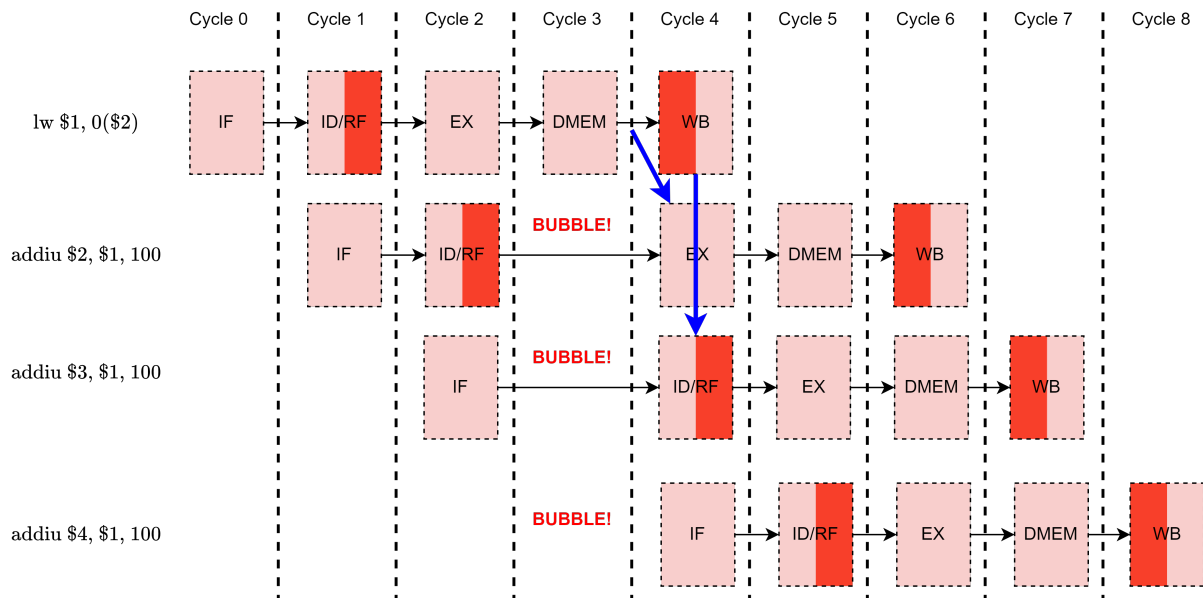
We can attempt to solve this issue using the compiler (e.g reorder instructions to put at least one non-dependent instruction between the load and the use).

Forwarding Paths

Software Scheduling

2.3.3 Control Hazard

Control Hazard	**Definition 2.3.4**

A stall created by the delay between getting the result of some branch/jump and fetching the next instruction using that data.

Instruction fetch (without branch prediction) requires the conditional branch result to be known. Hence the number of stages between instruction fetch and when the branch condition is determined is the size of the stall resulting from a conditional branch.

- This is also true for jumps/unconditional branches where the address is provided by some register and arithmetic (e.g jump with offset)
- Branch prediction can be done dynamically (in hardware) or statically (specific branch likely, branch unlikely instructions used by compiler).
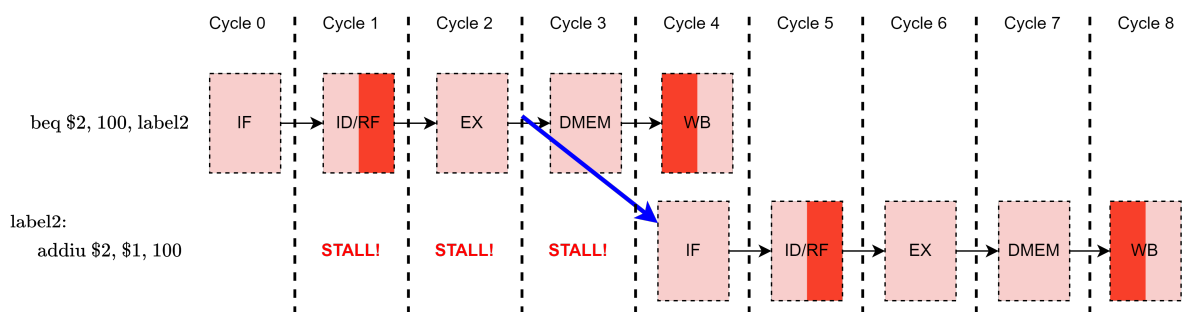
Early Branch Determination

To decrease the number of cycles stalled in a branch, we can move the branch determination to earlier in the pipeline.

- Instruction decode determines branch.
- Still a one cycle delay in the MIPS example pipeline above.

2.4 Simultaneous Multithreading

We can eliminate stalls by interleaving the instructions of independent programs.

- Maintain two program counters for two programs, two sets of registers. Alternate between instructions from each program.
- No dependencies between adjacent stages, less forwarding, less complex instruction decode and control required.
- Each program sees half the clock frequency.

2.5 Pipelining Roundup

Pipelining offers increased throughput without much added hardware complexity by allowing execution stages to run in parallel as a pipeline.

- Simple 5 stage pipeline can run at $5 \rightarrow 9$GHz
- Limited by critical path through slowest pipeline stage
- Clock period is $330ps \approx 10$ gate delays at $3GHz$ ($3 \rightarrow 5$ FO4 for latches, $5 \rightarrow 8$ FO4 for work).
- Memory access needs to be done in $5 \rightarrow 8$ FO4 delays (large constraint).

FO4 Delays	*Extra Fun!* 2.5.1
The gate delay of a component with a fan-out (gate inputs driven by a gate's output) of 4.	

Chapter 3

Caches

3.1 Why Caches

The difference between cycle time (time of a stage in a pipeline) and memory access time has continually increased.

Size	Access Time	Storage/Memory	Managed By	Transfer Unit
$100Bs$	$< 1ns$	Registers	programmer/compiler	$1 - 16Bs$
$10KBs$	$1ns$	Cache (SRAM) — L1, L2, L3	cache controller	$8 - 128Bs$
$100KBs$	$10ns$			
GBs	$100ns$ / $300ns$	Main Memory (DRAM)	Operating System	$4 - 8KBs$
TBs	$10ms$	Secondary Storage (Disk, SSD, Flash)	user/operator	MBs
$unbounded$	$seconds$ / $minutes$	Backup Storage (Tape)		

3.2 Locality

Programs typically access only a small part of their address space during a short time period.

Temporal Locality	Definition 3.2.1
Locality in time. The same location referenced is often referenced multiple times.	

Temporal Locality	Definition 3.2.2
Locality in space. Locations near an accessed location tend to be referenced soon.	

Most modern architectures are reliant on locality to determine when and what locations should be cached.

- Cache is a scarce resource.
- Cache misses are expensive.

3.3 Cache Types

3.3.1 Directly Mapped Cache

Associativity Conflicts	Definition 3.3.1

Where two or more locations are mapped to the same cache line/set of cache lines, and repeatedly replace each other.

```
/* Example with arrays, assume cache line is 256 bytes
 * and both arrays start at same cache index
 */

int array_a[64];
int array_b[64];

int some_function() {
    int sum = 0;
    for (int i = 0; i < 64; i++) {
        r +=
        array_a[i]    /* array_a moved into cache line */
        + array_b[i]; /* array_b evicts array_a and replaces */
    }
    return sum;
}
```
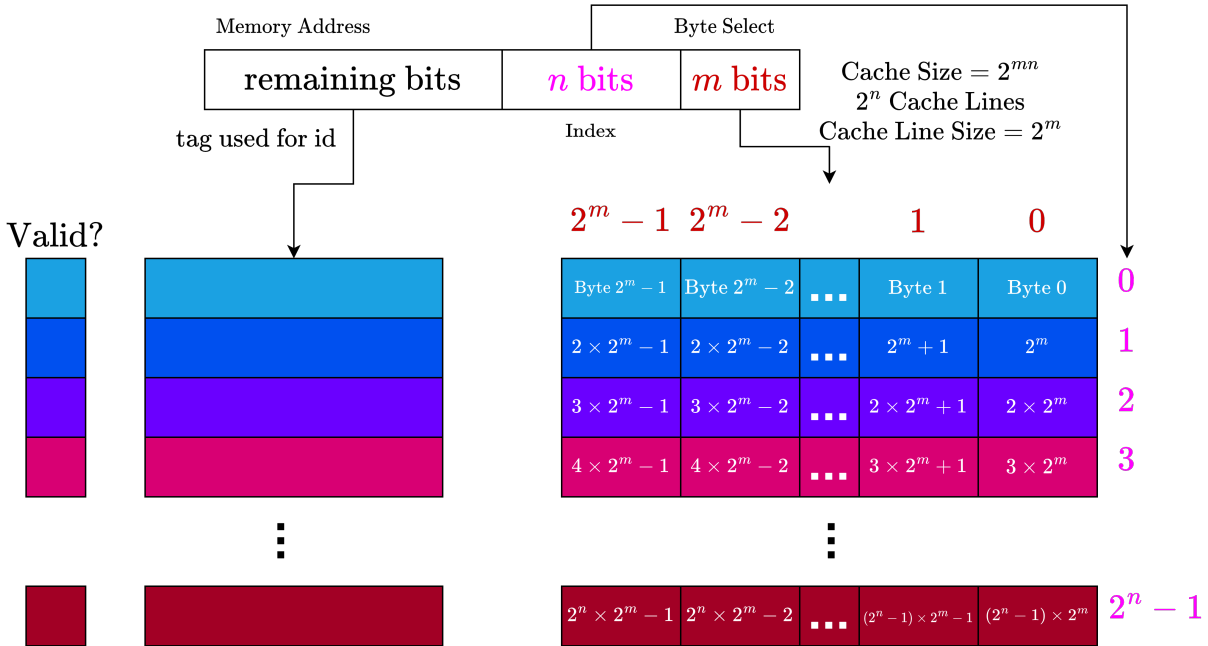

- Index and byte select used to find entry. Then tag compared to determine hit/miss.
- We can see a pattern in memory of where locations can be cached based on the index.

- Block/line received before the hit/miss is known (recover later if miss).

Memory   **Associativity Conflicts**

2^m Bytes

Simplicity	Simple indexing of cache & compare to determine hit/miss.
Fast Lookup	Only one location where a cached value may be.

Associativity Conflicts	As location can only be cached in one place, associativity conflicts are common.

3.3.2 Two Way Associative

Combine two directly mapped caches, and only cache a given location in one.
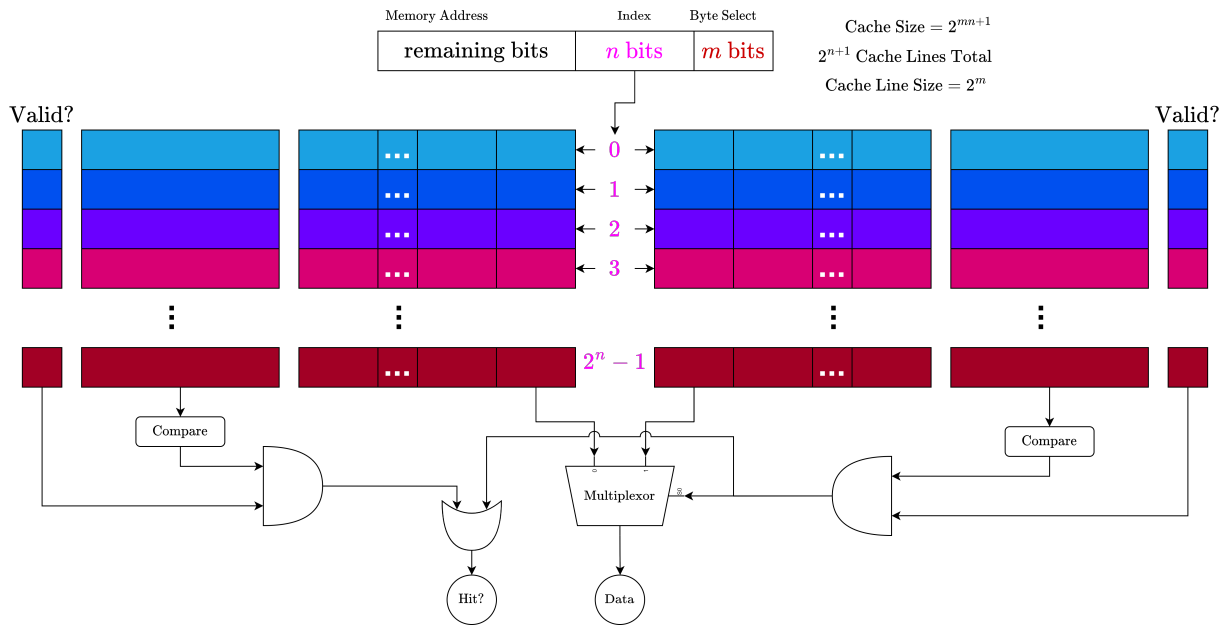
- Both caches searched in parallel.
- Only one hit possible, this is selected from result of both caches (selection is in the critical path)
- Cache block/line is available after the hit/miss is determined.

Fewer Assoc Conflicts	Any location can now select two different locations in the cache, hence two addresses with the same index can both be cached.

Multiplexer Delay	A multiplexer is added in the critical path
Complexity	Requires more comparators, and more complexity in placement & replacement.

3.3.3   N Way Associative & Block Placement

A generalisation of the directly mapped and two way associative caches. Block placement is restricted by the cache's associativity.

- Increasing associativity reduces associativity conflicts $\Rightarrow$ better hit rate (with diminishing returns)
- Greater overhead in terms of multiplexers in the critical path and the hardware complexity
- Fully associative cache can place any location in any cache location, and uses parallel search of tag (index is 0 bits) to find entry
- More associative $\Rightarrow$ less sensitivity to storage layout

---

Intel Pentium 4 Level 1 Cache (pre-prescott)   *Extra Fun!* **3.3.1**

Capacity:	$8KB$	**Block/Line Size:**	$64B$ so $^{8K}/_{64} = 128$ blocks
Ways/Associativity:	4	**Sets:**	32 (128 blocks, but 4 way $\Rightarrow$ $^{128}/_4$)
Index:	5 bits	**Tag:**	21 bits

Resulting access time is 2 cycles ($6ns$ at $3GHz$), with cache/memory being dual ported (load and store).

3.4   Block Identification

Index and tag identify a block.

- Increasing associativity decreases index size, increases tag size.
- Increasing block size decreases index size.

3.5   Block Replacement

When introducing a new location to the cache & possible locations are full.

- No choice in directly mapped.
- n choices for n-way associative.

The least recently used (LRU) evicts the oldest cache entry

- In practice only a marginal advantage over random eviction.
- Can be pathologically bad (e.g a loop accessing many locations may evict the first just before restarting the loop & accessing again).

3.6   Write Strategy

Write Through   **Definition 3.6.1**

On a cache hit, write to cache, and to the block in lower-level memory.

- Combined with write buffers to prevent a wait on memory
- Can always discard cached data, the most up to date si always present in memory
- Only requires a valid bit (cache control metadata)

Simpler	Cache management is simpler as the most up to date data is always in memory also.
Sharing	Next level of cache/potentially memory has the most up to date data.

Write Back   **Definition 3.6.2**

On a cache hit, only write back when evicting from cache.

- Track write backs with a dirty bit
- Absorb cost of repeated writes
- Cannot discard cache, when evicting it must be written back to memory
- Cache entries require both valid and dirty bits (cache control metadata)

Bandwidth	Memory is often overwritten several times, with write-back this will only require a memory write back when the cache entry is evicted.
Tolerance	Fewer memory accesses can result in a better tolerance to longer-latency memory (cheaper).

Write Allocate   **Definition 3.6.3**

When a cache miss occurs on write, allocate a new cache line and write to it.

- A read miss is required to fill in the rest of the cache line.
- As only *part* of the line is valid, a valid bit is required per word.

Write Non-Allocate / Write Around   **Definition 3.6.4**

When a cache miss occurs on write, send the data to memory / lower cache level (do not allocate a cache line).

Neither avoid the cache-coherence problem (inconsistent values for locations cached on multiple cores/processors).

3.7   Miss Rate Reduction Using Hardware

$$\text{Average Memory Access Time (AMAT)} = \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty}$$

In order to reduce AMAT:

- Reduce Miss Rate.
- Reduce Miss Penalty.
- Reduce time to hit cache.

3.7.1   Reducing Misses

Compulsory	First access so not in cache, also called a *cold start miss* or *first reference miss*.
Capacity	Cache cannot contain all blocks needed during the execution of a program. A capacity miss occurs when a block discarded due to capacity is later retrieved.
Conflict	Where the block placement strategy results in blocks being discarded as too many are mapped to a set (associativity conflicts). Also called *collision misses* or *interference misses*.

Compulsory	**Capacity**	**Conflict**
Infinite Cache	Fully associative, finite cache	n-way associative, finite cache

Coherence Miss   *Extra Fun!* **3.7.1**

A miss caused by cache coherence protocols. For example another core or an I/O device may invalidate a cache entry.

3.7.2 Increase Block Size

Spatial Locality Larger block means more locations are speculatively cached.
Cold Misses As more speculatively cached, fewer cold misses (cached speculatively before first access).

Wasted Space More of the cache is wasted on speculatively cached but unused data.
Conflicts larger Blocks $\Rightarrow$ Fewer lines $\Rightarrow$ increase capacity conflicts as there is potentially more contention over lines.
Loading Larger blocks may take longer to load (increased miss penalty) or require a wider bus (expensive hardware).

3.7.3 Increase Associativity

Fewer Associativity Conflicts More ways $\Rightarrow$ more ways to not conflict. This reduces the miss rate.

Cycle Time Comparators in the critical path

3.7.4 Victim Cache

The main cache is a large directly mapped cache. A victim cache is fully associative and smaller, and contains data discarded from the main cache.

- Checked in parallel.
- Rarely used for L1 cache, but often used for last-level caches.

This is an example of combining two strategies to avoid both's worst case behaviour.

Competitive Algorithms	Extra Fun! 3.7.2

Given two strategies, combining to create a good composite strategy.

- Ski Rental Problem (Combining the renting & buying strategies)
- Spinlocks vs context-switching (e.g spin before blocking?)
- Paging (replacement & eviction)
-

3.7.5 Skewed-Associative Caches

Given a two way set-associative cache.

- Use a different *hash function* for each cache (e.g one using regular (tag and index), other xors bits of index & tag and reorders)
- Associativity conflicts may be present in one cache, but not the other.

Address A Address B Address C

Associativity Conflict in Cache No Conflict

Valid? Valid?

f g

Different Hashes

Compare Compare

Multiplexor

Hit? Data

Some Conflicts	Its very difficult to write a program free of all conflicts.
More Decoders	Need an address decoder per way/cache rather than a single for all.
Hash Latency	Complex hash increases latency (it is in the critical path).
LRU	Difficult to implement LRU eviction policy (though this is not necessarily a good policy).

3.7.6 Hardware Prefetching

When a cache miss occurs, fetch the data (as required), but also pre-buffer the next block in a *stream buffer*.

- Prefetched blocks are placed in the cache (would pollute and potentially evict blocks to be used).
- Stream buffer checked in parallel with cache.
- Can add several prefetch buffers (multi-way stream buffer) to prefetch up w-way, fetch to w blocks ahead.
- Used in most high performance processors.

Sequential Access	Can avoid misses when traversing arrays.
Cache Untouched	Can use any type of cache design with this - an addition.

Memory Bandwidth	Need extra bandwidth to transfer block selected (cache miss) and block for pre-fetch.

Decoupled Access-Execute   *Extra Fun!* **3.7.3**

Decouple the processor into an access and execution sides.

- Access side fetches data to provide to the execute side.

- Execute side takes data from access and runs arithmetic instructions on it.

- Access side can be far ahead of execute, streaming the required data to it at close to memory bandwidth.

3.8 Miss Rate Reduction Using Software

3.8.1 Software Prefetching

Many modern processors provide prefetching instructions.

- Rarely needed - hardware prefetching is good!

- Useful on simpler processors with less or no hardware prefetching.

- Care required to prevent unwanted side effects.

- Prefetch instructions may target addresses that result in a page fault/protection violation (here they silently fail).

3.8.2 Reducing Instruction Cache Misses

Associativity conflicts can occur in the instruction cache.

- We want to avoid hot loops calling functions who's code have an associativity conflict with eachother.

- By using the caller graph, with each loop labelled, we can determine how to pack subroutines into the program binary to avoid associativity conflicts.

- Needs to consider the entire program, and the layout of all subroutines so must be done at link-time.

3.8.3 Storage Layout & Iteration Space Transformations

Merging Arrays	Improve spatial locality by merging two arrays into a single array of compou elements (i.e a zip). (Struct of Arrays vs Array of Structs)
Multidimensional Array Permutation	Match array layout to traversal order.
Loop Interchange	Change nesting of loops to access data in order stored in memory.
Loop Fusion	Combine independent loops that have the looping behaviour (e.g bounds) an overlapping variables. Sometimes this can then enable *Array Contraction*, whe some array can be replaced by a scalar value.
Blocking	Improve temporal locality by accessing cache line sized blocks of data repeated instead of accessing columns or rows.

A traversal order for blocks.

- Split blocks into four.
- Traverse four blocks in Z shape, recursively.
- A texture caching layout used in some GPUs.

```haskell
data QuadTree a = Single a |
  Quad {
    topLeft :: QuadTree a, topRight :: QuadTree a,
    bottomLeft :: QuadTree a, bottomRight :: QuadTree a
  }

{-
 A----->B
       |
 -------
 |
 C----->D
-}

morton :: (a -> b) -> (b -> b -> b -> b -> b) -> QuadTree a -> b
morton fun collect (Quad {tL, tR, bL, br})
  = collect (morton' tL) (morton' tR) (morton' bL) (morton' bR)
  where
    morton' :: QuadTree a -> b
    morton' = morton fun collect
morton fun _ (Single s) = fun s
```

3.9 Miss Penalty Reduction

Chapter 3 - Part 3: Reducing Miss Penalty

3.9.1 Write Buffers
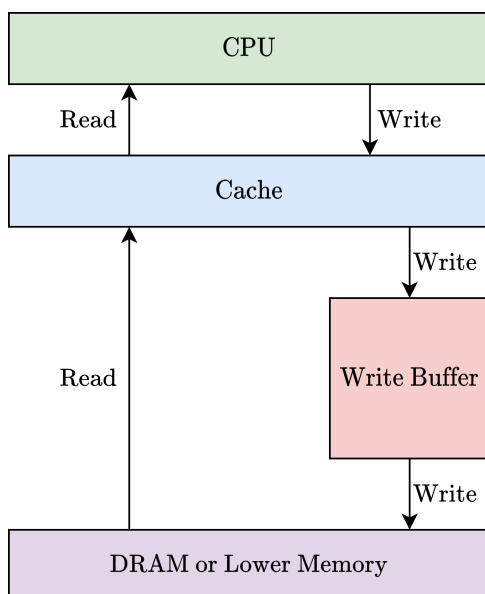
- RAW conflict on main memory reads when the location has a write in the write buffer.
- One solution is to wait until the write buffer is empty before reading, however this increases the read cache miss penalty
- A better solution is to check the write buffer on every memory read, if present in the buffer, take the value from there, else go to memory.

With write back we can reduce the stall for a read-miss that evicts a dirty cache line by:

1. Read miss on cache, evict dirty block.
2. Write dirty block to write buffer (fast).
3. Start read, CPU can resume/end stall when read is complete.
4. After read, write from write buffer to memory.

A cache is structured in terms of lines, hence the eviction of a cache entry means an entire line must be written

back to memory.

- Larger memory writes require more time, or more expensive/wider buses.
- The write buffer needs to be large enough to store multiple lines being evicted. Small write buffer will lead to stalls when full.

Coalescing Write Buffers	Adjacent writes are merged into a single entry in the write buffer. This is especially important in write-through caches.
Dependency Checks	Use comparators to check load addresses against pending stores. On a match a dependency is present, so the load must be stalled (other instructions can run).
Load Forwarding	If a store and load match address, forward the data to the load.

3.9.2 Early Restart

Sectored Cache Lines	Definition 3.9.1

A cache line can be divided into sectors.

- Each has its own validity bit, potentially dirty bit also.
- Cache allocated in units of cache lines
- Data delivered to cache in units of sectors.
- Sectors can be fetched in any order, potentially even remaining invalid until requested.

During a read-miss induced stall, the processor can restart as soon as the requested word arrives.
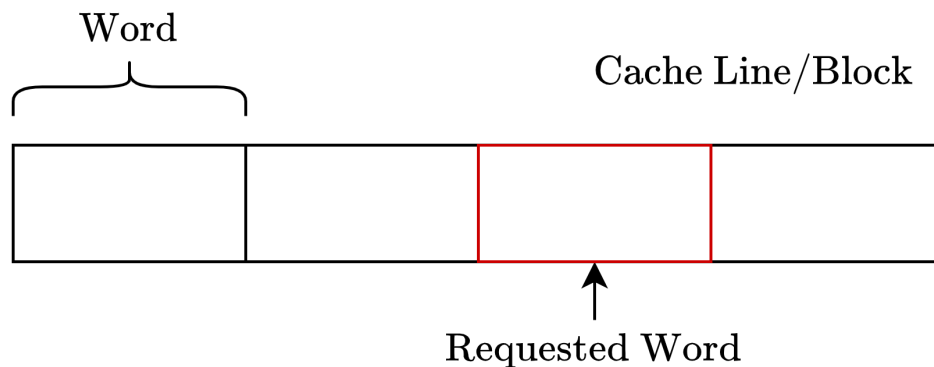
A cache line consists of many words, when loading a line into cache following a read-miss, we can restart the CPU as soon as the requested word is present.

Early Restart	As soon as the requested word arrives, send to the CPU and end stall/restart.
Critical Word First	Request the word on which the read-miss occurred from memory first.

Complexity	Sectoring and ordering reads increase hardware complexity. Must be careful for edge cases (e.g read-miss on cache entry that is currently in the process of being loaded).

3.9.3 Non-Blocking Cache

Non-Blocking / Lockup-free Cache	Definition 3.9.2

Allows data cache to continue to supply hits for other locations, during a cache miss.

- Requires full/empty bits on registers, and out-of-order execution.
- Requires multi-bank memories

Hit Under Miss	Effective miss penalty reduced as useful work is completed during a miss.
miss under miss	Misses are overlapped to reduce effective penalty.

Cache Controller Complexity	Needs to support multiple outstanding memory accesses to support *miss-under-miss*. Requires extra hardware (e.g multiple memory banks), and complexities of out of order execution.
Fences	Hit under miss allows for load to be services out of order, hence a fence/barrier instruction must be available to prevent this when required.

With In-Order pipeline processors, it is possible to implement some of this functionality by effectively making memory accesses out of order only.

- Freeze pipeline in Mem stage, but continue the rest.
- Use full/empty bits on registers, and a MSHR (Miss Status/Handle Registers) queue where each entry tracks the status of an outstanding memory request, a register may be marked as *empty* from a load, it will only stall if still empty when another instruction uses the register, at the decode stage.
- This is a popular approach with in-order superscalar processors.

3.9.4 Multiple Cache Levels

$$AMAT = \text{Hit Time}_{L1} + \text{Local Miss Rate}_{L1} \times \text{Miss Penalty}_{L1}$$

$$\text{Miss Penalty}_{L1} = \text{Hit Time}_{L2} + \text{Local Miss Rate}_{L2} \times \text{Miss Penalty}_{L2}$$

Can continue recursively for $L3, L4$ etc...

Local Miss Rate Definition 3.9.3
$$\frac{\text{Misses for this Cache}}{\text{Accesses to this Cache}}$$ Misses in a given cache, divided by the total number of memory accesses to the cache. • Relevant as a cache may not be accessed often, if the cache a level higher has a high hit rate.

Global Miss Rate Definition 3.9.4
$$\frac{\text{Misses for this Cache}}{\text{Total Memory Accesses}}$$ Misses in a given cache, divided by the total number of memory accesses.

Multilevel Inclusion Definition 3.9.5
Where lower caches contain all entries in the higher caches. $$d \in L1 \Rightarrow d \in L2 \Rightarrow d \in L3$$

- Inclusion strategy affects placement of data, and hence cache controls on coherence, search (should lower caches be checked?) etc.
- Can use the L2 cache to filter coherence protocol invalidations. If not in L2, no need to check L1, and hence no need to impact bandwidth to L1.
- L3 / LLC (Last Level Cache) often managed as a victim cache, only allocated when data is displaced from L2.

3.10 Hit Time Reduction

Chapter 4 - Part 4: Hit Time Reduction and Address Translation

3.10.1 Parallel Cache Access

When attempting to issue multiple instructions per cycle, parallel accesses to the cache are required.

Multiple Banks	Divide the cache into several banks, with addresses mapped to banks (e.g using low order bits, or a hash function). Accesses to different banks can occur in parallel.
Duplicate the Cache	By using multiple copies of the same cache, each can be accessed separately.
Multi-ported RAM	Add another wordline per row, and another bitline per column to allow multiple accesses to the RAM at the same time, cache uses this multiported RAM. This is effectively duplicating the cache, but sharing the flipflops between caches.

3.10.2 Address Translation

PIPT diagram: CPU → VA → TB (Translation) → PA → Cache L1 (PA Tags) → PA → Memory/Lower Cache. Physically-Indexed Physically-Tagged (PIPT). Some OOO Processors attempt to hide translation latency with reordering.

VIVT diagram: CPU → VA → Cache L1 (VA Tags) → VA → TB (Translation) → PA → Memory/Lower Cache. Hit does not require translation. Virtually-Indexed Virtually-Tagged (PIPT). When context switching, the cache must be flushed as it uses virtual addresses for tags.

VIPT diagram: CPU → VA → Cache L1 (PA Tags) and CPU → VA → TB → PA → Memory/Lower Cache. Cache Index determined by Virtual Address. Cache tag comes from Physical Address. Virtually-Indexed Physically-Tagged (VIPT). Cache can be kept when context switching, but latency of translation hidden with parallelism.

Homonyms   **Definition 3.10.1**

Same sound, different meaning

The same virtual address can point to different physical locations, in the context of different processes.

- A virtually indexed cache, as the tags are compared to determine hit/miss, we must flush the cache.
- In a TLB, where the cache is necessarily virtually tagged & indexed, a flush must occur, unless some process identifier is included in the tag (e.g ASID - Address Space ID).

Synonyms   **Definition 3.10.2**

Same meaning, different sound

- Multiple virtual addresses point to the same physical.
- In a virtually indexed cache multiple addresses may map to the same page.
- Updates to one cached copy must be reflected in the others (shared pages)

Faster Cache Hits by Avoiding Translation

Increase associativity to allow for a larger L1 cache without increase page offset size

Size of cache limited by size of page offset

Use some untranslated bits from the page number in the index

if A and B are different virtual addresses, that map to the same frame

The the untranslated page number bits used in the cache index must be the same for both these addresses, so they both use the same location in the cache, and hence any reads/write are reflected for both A and B

Linux MMap	*Extra Fun!* 3.10.1

`mmap` in linux does not specify the address in order to allow the operating system to determine this. This allows the OS to perform system specific tricks such as ensuring the page number bits in the cache index are identical for virtual addresses mapped to the same frame.

The L2 cache makes use of physical addresses, and hence relies on the virtual-to-physical mapping provided by the OS.

- The OS may choose mappings that result in associativity conflicts in the L2 cache.
- This means that different instances of the same program, with the same computation can have different performance, as the OS may assign frames such that an associativity conflict occurs.

Operating systems can also use different page sizes.

- large pages require smaller page tables, and use fewer TLB entries, which increases the hit rate for the TLB.
- Requires support in hardware for multiple page sizes.

Chapter 4

Dynamic Scheduling

4.1 Bypassing Stalls

The basic concept behind out of order scheduling is that instructions behind a stall can be allowed to continue provided data dependence/hazards allow.

- When an instruction stalls (e.g cache miss or forwarding not possible) save the state of that instruction.
- Instructions are issued in order, have dependencies analysed and can then be executed out of order.
- When operands are available allow execution pof the stalled instruction to continue.

Read After Write / True Dependence   **Definition 4.1.1**

```
add $3, $2, $1  # $3 = $2 + $1  (Write $3)
sub $4, $3, $6  # $4 = $3 - $6  (Read  $3) (needs previous instruction's value)
```

The output of one instruction is required as the input to another.

Write After Read / Anti Dependence   **Definition 4.1.2**

```
sub $4, $3, $6  # $4 = $3 - $6  (Read  $3) (use $3 before the next instruction overwrites)
add $3, $2, $1  # $3 = $2 + $1  (Write $3)
```

Some instruction will overwrite an input to a preceding instruction.

Write after Write / Output Dependence   **Definition 4.1.3**

```
add   $3, $2, $1  # $3 = $2 + $1  (Write $3)
sub   $3, $4, $6  # $3 = $4 - $6  (Write $3)
addiu $7, $3, 100 # $7 = $3 + 100 (Read $3)
```

The writes have a dependency as they write to the same location, the correct value must be present in the location for subsequent reads.

4.2 Tomasulo's Algorithm

An out of order execution algorithm used to dynamically rename registers to bypass the limited number of floating-point registers in the IBM architecture specification, and allow faster computation on the IBM 360/91.

- Each registers contains a tag. (null means the value is present, otherwise it is the identifier of the unit the result will come from)
- By adding tags register renaming (simple) is achieved

- A common data bus is used to broadcast the result of an operation, with its tag (unit it came from)

```python
"""Super abbreviated pseudocode for the IBM360/91 Out of Order Execution """
class IBM36091:
    def issue_instruction(instruction: Instr):
        # for each argument, check if register value is present or waiting.

        unit: FunUnitId = get_unit_from_opcode(get_opcode(instruction))
        dest: Register = get_dest_register(instruction)
        operands: List[Register] = get_operands(instruction)

        # overwrite destination with new unit to take result from (all subsequent instructions use result
        dest.set_tag(unit)

        # Get arguments (some from registers, some )
        args: List[ArgType] = []
        for op in operands:
            if (tag := op.get_tag()) is not None:
                args.append(WaitFor(tag))
            else:
                args.append(Value(op.get_register_value()))

    class Unit:
        def broadcast(data):
            # Broadcast data to registers and other functional units via the common data bus
            common_data_bus.broadcast(self.unit_id, data)

        def recieve(unit_id, data):
            # Given some data broadcast determine if it is needed, and if any instructions can be execute
            for instr in self.reservation_station:
                # Check if an instruction is waiting for the tag
                instr.take_args(unit_id, data)
                if instr.is_ready():
                    instr.queue_execute()

    class Register:
        def recieve(unit_id, data):
            # Check if the data broadcast is for the register.
            if self.tag == unit_id:
                self.tag = None
                self.value = data
```

Complexity	Led to delays in design, hardware overhead to overcome an ISA issue.
Limited by CDB	CBD must go through all functional units, and only one instruction can write to bus per cycle.
No Precise Interrupts	As instructions are executed out of order, we cannot clearly define a point in the *in-order* program text where the processor is at at any given time.

It is possible to overlap loop iterations:

- (Effectively) Register renaming allows for different physical destinations (e.g ignore register and straight to functional unit).

- Reservation stations can buffer old values to avoid write after read / anti dependence stalls.

Chapter 2 - Part 2: Speculation

4.3 Precise Interrupts

In order to use precise interrupts we need a consistent state.

- All instructions up to some point have committed changes to machine state (registers & memory).
- No instructions past have committed.
- Hence on an interrupt (e.g page fault, syscall) we can easily save state, and restart where the interrupt suspended execution of a program.
- This is also important for branches (need to undo prevent committing work executed speculatively)

Hence we want to make a *speculative tomasulo algorithm*

1. Issue/Dispatch (Get instruction from buffer of fetched instructions, send operands & reorder buffer number to destination)
2. Execution (Out of order execution of issued instructions)
3. Write Back (in order to common data bus and waiting functional units)
4. Commit (Update register with reorder result, reorder buffer takes completed instructions, puts in issue order and updates state)

This requires several additions

- Commit unit to manage reorder buffer
- Issue side registers for execution
- Commit side registers for the committed results
- Ability to flush the reorder buffer on a branch mispredict

4.4 Store Buffering

Stores are an issue as they cannot be completed until committed, but succeeding loads can be executed straight away.

- We could stall all preceding loads until the store is complete
- We can buffer uncommitted stores, associated with addresses, and check these for any load (to get the nearest hit, or on miss go to memory). Loads must be stalled until all possibly aliasing store addresses are resolved

Loads and stores use computed addresses (not always known at issue time)

- Can speculate, and forward a store's result to a load
- Must recover when the computed address is not the speculated

Hence we can add a *forwarding predictor* to determine if a store should be forwarded to some load behind it in the pipeline.

Dependence Prediction   *Extra Fun!* 4.4.1

More can be read about predicting the dependence of a load on another store instruction https://jilp.org/vol2/v2paper13.pdf.

4.5 Register Update Unit

An alternative to reservation stations and the reorder buffer.

- A single table of instructions after fetch, acting as a reservation station.
- Once the operands are found, the instruction can be issued (hence functional unit determined after operands, unlike in Tomasulo's)
- RUU entries are committed to update the commit side registers.

Monitors	In Tomasulo's every reservation station and reorder buffer entry needs to have a comparator and monitor the common data bus. With the RUU strategy, fewer comparators are required.
Tags	With RUU the tags are ROB entries. Furthermore the RUU is indexed by the tag.

4.6   Register Alias Tables

UNFINISHED!!!

Chapter 5

DRAM

UNFINISHED!!!

Chapter 6

Side Channels

Chapter 5 - Part 1: Sidechannel Vulnerabilities

Side Channel Attack   **Definition 6.0.1**

An exploit that attacks the implementation of an algorithm by observing the state of system it runs within. For example:

- Detecting what is cached through memory access times
- Power consumption, electromagnetic leaks and other physical effects

6.1 Exfiltration

Prime and Probe   **Definition 6.1.1**

Detect eviction of the attacker's working set, that is caused by the victim.

1. Attacker primes the cache by filling some sets with its own lines.
2. Victim executes, once finished the attacker probes (timing its memory accesses) to see which of its lines were evicted.

If a line has been evicted, then the victim accessed an address mapping to the same set.

- Requires the attacker to be able to force the start of victim execution.

Evict and Time   **Definition 6.1.2**

Detecting cache usage by the victim, by monitoring its performance after altering the cache.

1. Attacker causes victim to execute, and to preload the cache with its working set.
2. Attacker evicts a specific line from the cache.
3. Victim executes, attacker monitors execution time.

By repeating this process for many lines, the attacker can see where the time to execute is lower, and hence which lines the victim was using.

- Requires the attacker to be able to force the start of victim execution.

6.2 Shared State

On a modern CPU a large amount of state is shared between cores.

- L2 and lower Cache
- core interconnects/buses

If two threads run on thew same core (as with SMT, or HyperThreading):

- Instruction and data caches, as well as TLB
- Branch predictor
- Prefetcher
- Rename Registers
- Dispatch Ports

6.3 Triggering Victim Execution

System Call	If the operating system can invoke the victim, or the operating system itself is the victim.
Lock Release	A victim may be waiting on a held lock (e.g a lock on a file in the filesystem)
Call it	If the victim is contained in the same address space, call it as a function.

The latter point (*"call it"*) applies as the attacker's code may be running in the same process as the victim, under some runtime system (e.g the JVM, or a browser's javascript engine).

> **Language Based Security**   **Definition 6.3.1**
>
> The runtime system bundled with and provided by the language enforces separation between threads.
>
> - i.e prevents pointer arithmetic, checks array index bounds

6.4 Side Channels in Speculative Execution

Speculative execution of instructions can impact the state of the cache.

- Can include instructions in code, that will be speculatively executed and as a result impact cache.
- e.g Bypassing array bounds checks using the CPU's speculative execution of the invalid instruction to modify cache.

This vulnerability is commonly known as *Meltdown & spectre* are is present in most modern dynamically scheduled processors.

> **Chapter 5 - Part 2: Attacking Other Processes and the OS**

6.5 Mitigation

ASLR	Address Space Layout Randomisation is where the operating system distributes user pages across memory. It can also randomly distribute the kernel pages locations within a process' address space.
KPTI	Kernel Address Space Isolation is where the virtual address mapping of kernel pages is changed every time the kernel is entered. This requires a reload of the TLB and a substantial performance penalty.

6.6 Spectre v2

Even with these mitigations, it is still possible by training the branch predictor to mispredict a call within the victim's code.

1. Train the branch predictor to predict jumps to your subroutine. (Need to consider design of the BTB)

2. Make the syscall to the kernel.

3. Kernel code starts running, hits a call that is mispredicted to the attacker's code.

4. Attacker's code speculatively executes before being flushed, but effect on cache, size channels is still present.

This is called *Spectre Variant 2*

There are several possible mitigations:

Block Microarchitecture and Cache Side Channels	Very difficult to impossible.
Reduce Accuracy of Probing	Need to add noise to timers, but this will now affect other applications that need precision timing.
Prevent Branch Predictor Poisoning	Can add an instruction to prevent branch prediction, but there is a performance cost.
Block Branch Predictor Contention	Keep separate perdictions for each thread protection domain (ring/protection level).

Another solution is *retpoline* - an indirect branch using a return instruction.

```
RP0:   call RP2                   ; Push address of
RP1:   int 3                      ; A breakpoint (will be speculatively executed, but is a mispredict)
RP2:   mov [rsp], <jump target>   ; overwrite the return address, now return will go to <jump target>
RP3:   ret                        ; jump to jump target
```

Chapter 7

Exploiting Parallelism

7.1 Static Scheduling

Dynamic scheduling (out of order scheduling) requires significant hardware complexity.

- Register Update Units, reorder buffers, registers backed by commit registers and associated with tags, instruction dependency checks.
- All of these take space on the die (not only does a larger chip necessitate fewer chips per wafer, but the yield is also decreased)
- Also requires more energy, results in more heat and hence lower thermal limits.
- The complexity of determining the number of instructions that can safely be issued in parallel is $O(n^2)$, which is achievable for small n, but can necessitate more stages between fetch and issue.

With static scheduling this complexity is removed from hardware, and moved to the compiler, with the ISA providing necessary mechanisms to express how instructions should be scheduled (e.g in parallel).

7.1.1 Software Pipelining
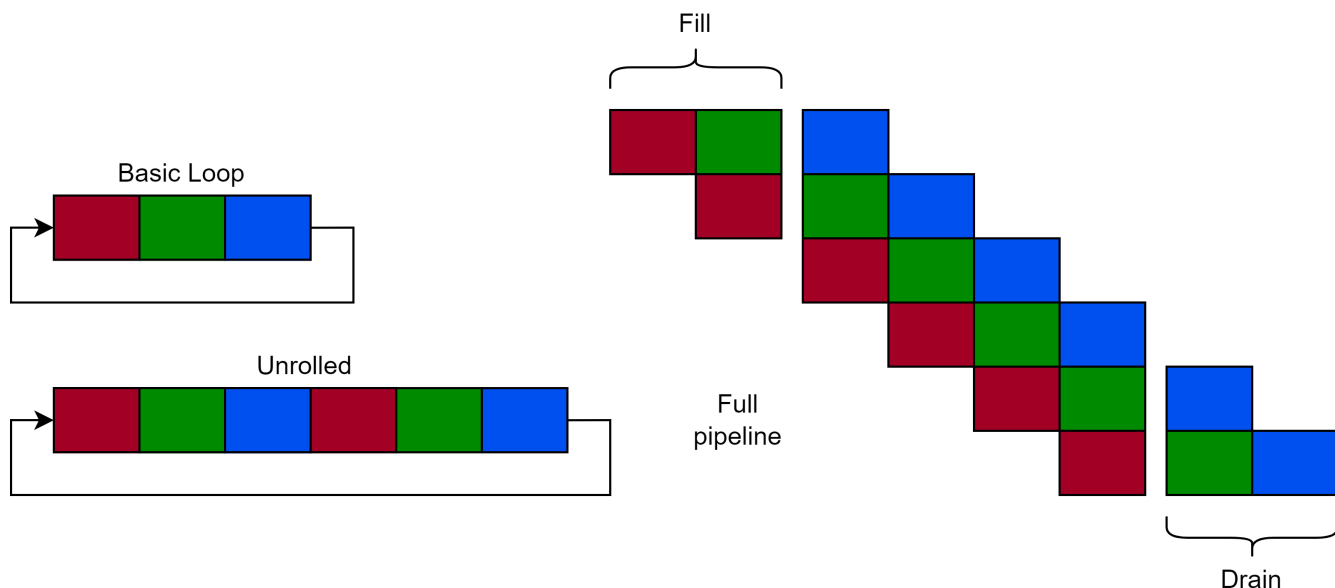
We can pipeline loop iterations, in the diagram above the basic loop and unrolled loop both execute the loop contents in order. By pipelining each of the 3 instructions in the loop body are run for 3 different in-order *iterations*.

- e.g iteration 3 is in red, while 2 is in green and 1 is in blue.
- Increases the load-use distance, so removes/reduces stall potential.

7.1.2 Very Lond Instruction Word

<div style="border:1px solid blue">

Very Lond Instruction Word (VLIW)   **Definition 7.1.1**

Each instruction contains encodings for multiple operations.

- All operations are independent and hence can be issued and executed in parallel.
- The compiler/programmer needs to extract dependencies, and work out which instructions can be issued & executed in parallel, rather than the hardware.
- Instructions become large, and where there is little parallelism to be extracted, majority of the instructions are mostly no-ops.
- Large instructions put pressure on memory access bandwidth.
- Often not binary compatible across generations (e.g number of functional units change, instruction size changes)

</div>

With software pipelining we can schedule instructions for different stages of the pipeline in parallel.

7.1.3 Explicitly Parallel Instruction Computing

<div style="border:1px solid blue">

Explicitly Parallel Instruction Computing (EPIC)   **Definition 7.1.2**

A term created by Intel & HP, considered to be the next generation of VLIW.

- Often used to refer to IA-64 (Itanium) processors.
- ISA exposes parallelism to the compiler.
- Binary compatible across generations/processor implementations.

</div>

In IA-64 instructions are encoded in bundles, each 128 bits wide:

- 5 bit template field encodes which instructions can be run in parallel in the bundle (where the ; ; / *stop*! is, after which the next set of parallel instructions begin)
- 3 instructions, each with 41 bits of length, this allows for large number of registers, large immediate operands.

Rotating Register File

- Registers 0 to 31 are always accessible.
- Registers 32 to 128 can rotate.

This allows for two main advantages:

Register Stack	By using a special register CFM (current frame pointer) to point to the set of registers used by a procedure. This allows many register arguments to be used for function calls, with results placed in registers in a stack like fashion.
Improved Software Pipelining	As the registers can be *rotated* we can pipeline loops more easily (register names remain the same, but values rotated for each loop iteration).

```
# Branch to L1, and rotate the register file by 1.
br.ctop <label> ;;
```

Predication

There are 64, single bit registers that can be used to determine whether an instruction is run.

- Branches can be eliminated in favour of predicated instructions (can hence avoid branch & branch mispredict costs)
- Can issue both sides of a branch in parallel & predicate both.
- Can easily move instructions across conditional branches.

Speculative Loads

- Compiler can specify speculative loads, and specify when loads will be used.
- Reduces cycles wasted to load-use stalls.
- Speculative loads do not fault, and hence can safely be used in code containing branches (e.g check if a pointer is null, can speculatively load before null check, but never use the value)
- An *advanced load* variant checks for aliasing stores (stores to same location)
- An *advanced load address table* tracks stores to addresses of *advanced loads*

```
# Speculatively load r1 = *a
ld.s r1 = [a]

# If the load for r1 faulted, go to some other branch, else NOP
chk.s r1

# r1 value made non-speculative and can now be used.
use = r1

# An advanced speculative load, monitors the address b for loads
ld.a r2 = [b]

# A store occurs to address b
st ??? [b]

# If aliasing has occurred for b, then re-load it
ld.c r2 = [b]

use = r2
```

When speculative loads are not fulfilled (i.e due to a fault - e.g page fault) NaT (Not a thing) is placed in the destination register.

- Speculatively loaded data can be consumed by other instructions before use.
- The NaT is propagated until checked.

UNFINISHED!!!

7.2 Multithreading

> **Chapter 7: Multithreading**

Time

Single Instruction Issue

A single instruction is issued per cycle

Chip Multiprocessor

Separate Cores on the same die running in parallel

Superscalar Issue

Processor analyses dependencies in a window of instructions to issue instructions in batches

Greater performance, however utilisation is reduced

Missed opportunities to issue

Fine Grained Multithreading

Round robin through n threads, hence any stall for less than n cycles can masked by the 4 cycle wait between instructions from a single thread.

intra-thread dependencies can still limit performance

Predicated Issue

Predication allows for fewer branches, and hence less branch related stalls, but many instructions will not be run/are wasted

Some predicated instructions will be discarded (predicate false)

Simultaneous Multithreading

All threads can have instructions issued at any time, and can

---

> **Fine Grained Multithreading (FGMT)**   **Definition 7.2.1**
>
> Each cycle one thread can issue instructions.
>
> - Typically round-robin through threads.
> - Hence can *hide* a stall of n cycles for n threads.

7.2.1 Simultaneous Multithreading

> **Simultaneous Multithreading (SMT)**   **Definition 7.2.2**
>
> Where instructions from several threads can be issued in any cycle.
>
> - Requires a more complex frontend (e.g to tag cache entries with which thread they are for, TLB needs to know page table per thread)
> - Instructions can be scheduled from any of the threads, maximizing utilisation.
> - Threads may contest resources (i.e two threads want to use the same functional units) resulting in reduced performance for a thread, conflicts in the cache.
> - Side channel attacks are an issue (threads share cache, scheduler needs to prevent one thread monopolising the CPU)
> - Some resources do not need to be replicated for each thread (e.g can have n times more logical registers, but not actual registers, one thread can use more than $1/n$ of the cache)

SMT can allow for memory-system parallelism to be exploited

- Lots of threads can have memory accesses *in-flight*.
- Can overlap data accesses with computation from other threads (e.g issue from thread B while A is stalled on a load-use)

34

UNFINISHED!!!

7.3 Vector Processing

	Chapter 8: Vector Instructions and SIMD

7.3.1 Arithmetic Intensity

Arithmetic Intensity

Sparse Matrix (SpMV)	Structured Grids (Stencils, PDEs)	Structured Grids (lattice Methods)	Sectral Methods (FFTs)	Dense Matrix (BLAS3)	N-Body (Particle Methods)

- Arithmetic intensity compares the ratio of arithmetic operations to memory operations.
- Hence we can determine if a program is limited by the rate of arithmetic operations, or by the memory bandwidth.
- FLoating point arithmetic is often used in this context, where arithmetic intensity is a measure of $FLOPs/Byte$
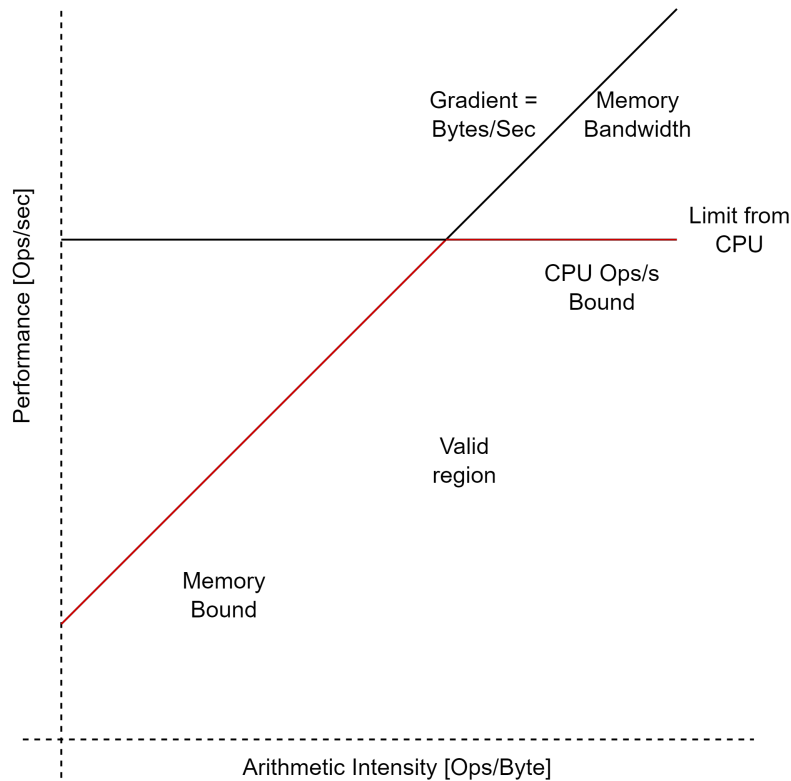
7.3.2 Vector Instruction Set Extensions

A vector extension for Intel's x86-64 architecture.

- 32 extended registers (ZMM0 $\rightarrow$ ZMM31), each is 512 bits wide.
- Can use registers to store 8 doubles, 16 floats, 32 shorts or 64 bytes.
- Instructions are executed in parallel in $4, 32, 16$ or 8 *lanes*.
- Predicate registers (k0 $\rightarrow$ k7 where k0 is always true), each predicate register holds up to 64 bits and hence each register can hold a predicate per *lane*

Subroutines available for use in a given language, with an implementation handled by the compiler.

- Often related to performance - e.g software prefetching.
- Used where the language cannot express some constraints / semantics - e.g vector instructions & C

Compiler intrinsics are provided for emitting specific vector instructions. For example with AVX12:

```
#include <immintrin.h>

// we can now use compiler instrinsic
res = _mm512_maskz_add_ps(k, a, b)

# Instead of assembly
VADDPS zmm1 {k1}{z}, xmm2, zmm3
```

Conditionals may be required in vectorised code, to allow this the predicate registers are used to determine the results of their corresponding lanes.

Zero Masking	z or {z} prefix	Inactive lanes produce a zero
masking	No prefix	Inactive lanes do not overwrite previous result.

The compiler must consider several issues in order to vectorise even simple code:

```
void add(int *a, int *b, int *c) {
    for (int i=0; i < N; i++)
        c[i] = a[i] + b[i]
}
```

Aliasing	In the above example a and b may overlap with c. We can inform the compiler these are separate using `restrict`. Otherwise the compiler may need to generate code that checks for aliasing, and only runs vectorised code if the arrays are distinct.
Size	The size may not be a size supported by vectorisation (here we can have 16 elements), may vector instructions may be needed (in a loop), and for a non-power-of-two extra loop code may be required.
Alignment	The vectorised load typically requires some alignment (e.g on a 32 byte boundary), extra code may need to check this alignment and sequentially execute part of the start/end of the loop.

Ultimately the compiler may not be able to vectorise a loop. We can use tools like OpenMP to tell/ask the compiler to vectorise code, even if it cannot assure this is safe.

```
// To inform the compiler vectorisation is safe (it may not vectorise)
void add(float *a, float *b, float *c) {
    // Ignore Vector DEPendencies / assume no loop-carried dependencies
    #pragma ivdep
        for (int i = 0; i <= N; i++)
            c[i] = a[i] + b[i]
}
```

```c
// To tell the compiler to vectorise the code (it may still not - e.g call
// non-vectorisable function from within loop told to bve vectorised)
void add(float *a, float *b, float *c) {
    // OpenMD make this loop SIMD
    #pragma omp simd
        for (int i = 0; i <= N; i++)
            c[i] = a[i] + b[i]
}


// We can also make functions vectorisable, hence this function can be called
// from a vectorised loop
#pragma omp declare simd
void add_element(float *a, float *b, float *c) {
    *c = *a + *b
}
```

If the compiler still will not vectorise, we can rely on intrinsics:

```c
void add(float *a, float *b, float *c) {
__m128* pSrc1 = (_m128*) a;
__m128* pSrc2 = (_m128*) b;
__m128* pDst = (_m128*) c;

// lengths are part of data types
for (int i = 0; i <= N / 4; i++, pSrc1++, pSrc2++, pDest++)
    *pDest = _mm_add_ps(*pSrc1, *pSrc2)
}
```

7.3.3 Single Instruction Multiple Thread

Each *lane* can be considered a thread, where all threads execute the same instructions synchronously/in lock-step.

- We can attempt to vectorise arbitrary control flow through the use of predicate registers.
- An outer loop can be vectorised, with the inner loops being iteration

```c
#pragma omp simd
for (int i = 0; i < N; i++) {
    // 1. set predicates for true, false branches.
    // 2. run predicated true branch
    // 2. run predicated (opposite) false branch (do not overwrite registers)
    if (...) {...} else {...}

    // continue running vectorised instruction until all predicates are false
    // do not overwrite when predicate is false
    for (...) {...}
    while (...) {...}

    // If the function can be vectorised, then a call to the function can be vectorised.
    function(...)
}
```

We can demonstrate this with a condition

```c
void add(float *a, float *b, float *c) {
    for (int i = 0; i <= N; i++)
        if (a[i] != 0.0)
            c[i] = a[i] + b[i]
}
```

Which can be compiled to the following assembly code:

```asm
add:
    xor eax, eax  # note this also zeros out all of rax
```

```
        vpxord zmm0, zmm0, zmm0   # zero-out zmm0
loop:
        # Load a[] into the zmm1 register
        vmovups zmm1, ZMMWORD PTR [a+rax*4]

        # Compare all the elements (each 4 bytes wide) and place results in k1 predicate register
        vcmpps k1, zmm1, zmm0, 4

        # Add b[] element wise to a[]
        vaddps zmm2, zmm1, ZMMWORD PTR [b+rax*4]

        # Conditionally move the result from zmm2 into c[]
        vmovups ZMMWORD PTR [c+rax*4]{k1}, zmm2

        # list iteration
        add rax, 16
        cmp rax, 1024
        jb loop

        # zero-out the upper bits of the zmm registerss
        vzeroupper

        ret
```

ARM Scalable Vector Extension (SVE)   *Extra Fun!* **7.3.1**

The arm ISA uses SVE for vector instructions:

- Has a First Fault Register (FFR) to allow for speculative loading of vectors (a page fault turns up in the FFR, program can continue), this helps with gather & scatter (indirection) where cache-miss induced stalls are more likely

- Hides the vector instruction width as an implementation detail (maximum 2048 bits)

Read more here.

Indirection is often used (e.g `array[other_array[i]]`) hence instructions are provided to load using a pointer in each lane.

- AVX512 has `vgatherdps`.
- Can result in many cache misses (and resulting transfers, evictions & allocation)

7.3.4 Vector Pipelining

In a vectorised loop, many iterations of some vectorised instructions may be required.

- We can software-pipeline this in much the same way as we have done with scalar instructions.
- Forwarding works heavily to our advantage here.

This can include breaking down vector instructions, for example if the floating point unit is only 8 wide, then we can pipeline 32 bit wide vector instructions into 4 blocks, and pipeline these.

7.3.5 Micro-Op Decomposition

The ISA may support wider vector instructions than it has ALU's for:

- Can split vector instruction into parts at decode, dynamically schedule parts and gather in the commit side.
- By breaking vector instructions down, halves can be dynamically scheduled - a delay in one for memory accesses does not stall entire instruction.

7.4   Graphics Processing Units

UNFINISHED!!!

Chapter 8

Parallel Programming

8.1 Motivation for Parallelism

Power is a critical constraint on the performance of a core executing a single thread.

Dynamic	Power consumed when signal is changed.
Static	Power consumed to power-up a gate.
Static Leakage	Charge is lost through quantum tunnelling (electrons skip across a gate). This increases exponentially as the gate size is reduced.

Denard Scaling   **Definition 8.1.1**

The dynamic power of a transistor decreases as the size of the transistor decreases.

- Smaller transistors use less power, and can be clocked at higher frequencies.
- Smaller transistors also help in increasing the hardware that can be placed on a die of a given size.

As transistor size has decreased, static leakage has come to dominate power usage, especially at high voltage (required for high clock rate).

- Need high voltage to move charge more quickly
- Higher voltage $\Rightarrow$ more leakage
- Chip must be kept within temperature limits to function

Rather than increasing clock rate (becoming very difficult) we can increase the parallelism in the chip & programs. This is generally much more energy efficient than increasing clock frequency.

There are several ways to mitigate power usage

Clock Gating	Turn functional units off when unused, deallocate part of the processor (e.g shut down part of the cache), potentially even entire cores (used with arm' big.little architecture)
Dynamic Voltage & Clock	Reduce performance by adjusting clock rate & voltage (e.g when battery low). When the processor is not the bottleneck (e.g blottlenecked by screen refresh rate, memory system) there i no need to *'speed ahead'*. This is very popular technique.
Spread Load	Run many cores at a low clock rate (parallelism).
Turbo Boost	When a single thread is running, can shut off other cores and increase the clock rate temporarily (boost clock rate) on the single core being used.

8.2 Shared Memory Parallelism

8.2.1 For Loops

> **OpenMP**   **Definition 8.2.1**
>
> OpenMP is a specification for language extensions to allow shared-memory parallelism.
>
> - Bindings exist for Fortran, C and C++ (with experimental implementations for Java and C#)
> - Allows the programmer to specify how a program should be parallelised
>
> ```
> // for example parallelism in for loops
> #pragma omp parallel for
> for (...) {...}
> ```

We can implement for loops in several ways, one is a self-scheduling loop:

```c
for (int i=0; i < N ; i++) {
    c[i] = a[i] + b[i]
}

int i;
if (myThreadID() == 0)
    i = 0;

// No thread can cross barrier till all have arrived.
barrier();

for(;;) {
    int local_i = FetchAndAdd(&i);

    if (local_i >= N)
        break;

    c[local_i] = a[local_i] + b[local_i]
}

barrier();
```

We can perform several optimisations here:

- Potentially avoid some barriers (depends on implementation of `fetchAndAdd` also)
- Work on chunks (if each thread is on a different core, then each has a different L1 cache & hence having each thread/core work on data with spatial locality is advantageous)
- Use cache affinity (previous for loop will have allocated entries in L1 caches of different cores, we can be smart about which cores which threads run on to take advantage of this)

8.2.2 Atomic Operations

Using locks is expensive (especially those that use syscalls). hence we can use atomic operations instead.

- Many languages provide mechanisms for using atomics (e.g `<atomic.h>`)
- Intrinsics can also be provided on a low level (e.g `__sync_fetch_and_add(p, inc)` in C)
- The instruction set mut provide some atomic mechanism, in intel this is the `LOCK` prefix, which ensures the operation occurs on a cache line held exclusively (no other cached copies)
- In a large system, atomics can cause contention (e.g many fetch & adds to the same location result in many threads waiting), in a network we can combine these atomic increments(e.g two fetch and increments become a fetch and add 2)

OpenMP supports several methods for this:

```
#pragma omp parallel for                                             \
    default(shared) private(i) All variables except i are shared between \
                               threads                                \
                                                                      \
    schedule(static, chunk)    Iterations of the loop can be distributed \
                               in equal sized blocks to each thread   \
                                                                      \
    reduction(+:result)        Perform a reduction on the variables that \
                               appear in the argument list (private copy \
                               created at the end, then all have their \
                               copies combined)
for (i = 0; i < N ; i++)
    result = result + (a[i] * b[i])
```

8.3   Distributed Memory Parallelism

> **Mesage Passing Interface (MPI)**   **Definition 8.3.1**
>
> A standard API for parallel programming using message passing.
>
> ```
> MPI_Init       // Initialise MPI
> MPI_Comm_size  // Get the number of processes
> MPI_Comm_rank  // Get this process
> MPI_Send       // Send a message
> MPI_Recv       // Receive a message
> MPI_Bcast      // Broadcast data fro the process with rank "root" to all other processes
> MPI_Reduce     // Combine values on all processes into a single value using an argument op (e.g sum)
> MPI_AllReduce  // MPI_Reduce and broadcast so every process has the reduced value
> MPI_Finalize   // Terminate MPI
> ```

The key idea with MPI is to use collective operations to write a collection of cooperating processess.

- General model to follow is that each process runs the same code (same control flow) and owns a share of the data (Single Program Multiple Data)

> **Stencil**   *Extra Fun!* **8.3.1**
>
> A stencil is a program that updates array elements (1d, 2d, 3d +) based on some fixed pattern (the stencil).
>
> - For example Conway's Game of Life uses a stencil to update a cell based on its neighbour's values.
> - This appears in image processing frequently (e.g blurring, image filtering)
> - Arises in convolutional neural networks, solving differential equations etc

We can demonstrate this with a stencil program.

With OpenMP

```
while(!converged) {
    #pragma omp parallel for \
        private(j)           \
        collapse(2)
    for(int j=0; j<M, ++j)
        for(int i=0; i<M, ++i)
            B[i][j] = 0.25 * (A[i-1][j]
                            + A[i+1][j]
                            + A[i][j+1]
                            + A[i][j-1])

    #pragma omp parallel for \
        private(j)           \
```

```
        collapse(2)
    for(int j=0; j<M, ++j)
        for(int i=0; i<M, ++i)
            A[i][j] = B[i][j]
}
```

With MPI

The following code is from this example and is written in fortran.

- Data is partitioned per process.
- Need to consider the *halo* (values just beyond the edges that must be read)

```
...
REAL, ALLOCATABLE A(:,:), B(:,:)
INTEGER req(4)
INTEGER status(MPI_STATUS_SIZE, 4)
...
! Compute number of processes and myrank
CALL MPI_COMM_SIZE(comm, p, ierr)
CALL MPI_COMM_RANK(comm, myrank, ierr)

! compute size of local block
m = n/p
IF (myrank.LT.(n-p*m)) THEN
    m = m+1
END IF

! Compute neighbors
IF (myrank.EQ.0) THEN
    left = MPI_PROC_NULL
ELSE
    left = myrank - 1
END IF
IF (myrank.EQ.p-1)THEN
    right = MPI_PROC_NULL
ELSE
    right = myrank+1
END IF

! Allocate local arrays
ALLOCATE (A(0:n+1,0:m+1), B(n,m))
...
! Main Loop
DO WHILE(.NOT.converged)
    ! compute
    DO i=1, n
        B(i,1)=0.25*(A(i-1,j)+A(i+1,j)+A(i,0)+A(i,2))
        B(i,m)=0.25*(A(i-1,m)+A(i+1,m)+A(i,m-1)+A(i,m+1))
    END DO

    ! Communicate
    CALL MPI_ISEND(B(1,1),   n, MPI_REAL, left,  tag, comm, req(1), ierr)
    CALL MPI_ISEND(B(1,m),   n, MPI_REAL, right, tag, comm, req(2), ierr)
    CALL MPI_IRECV(A(1,0),   n, MPI_REAL, left,  tag, comm, req(3), ierr)
    CALL MPI_IRECV(A(1,m+1), n, MPI_REAL, right, tag, comm, req(4), ierr)

    ! Compute interior
    DO j=2, m-1
        DO i=1, n
            B(i,j)=0.25*(A(i-1,j)+A(i+1,j)+A(i,j-1)+A(i,j+1))
```
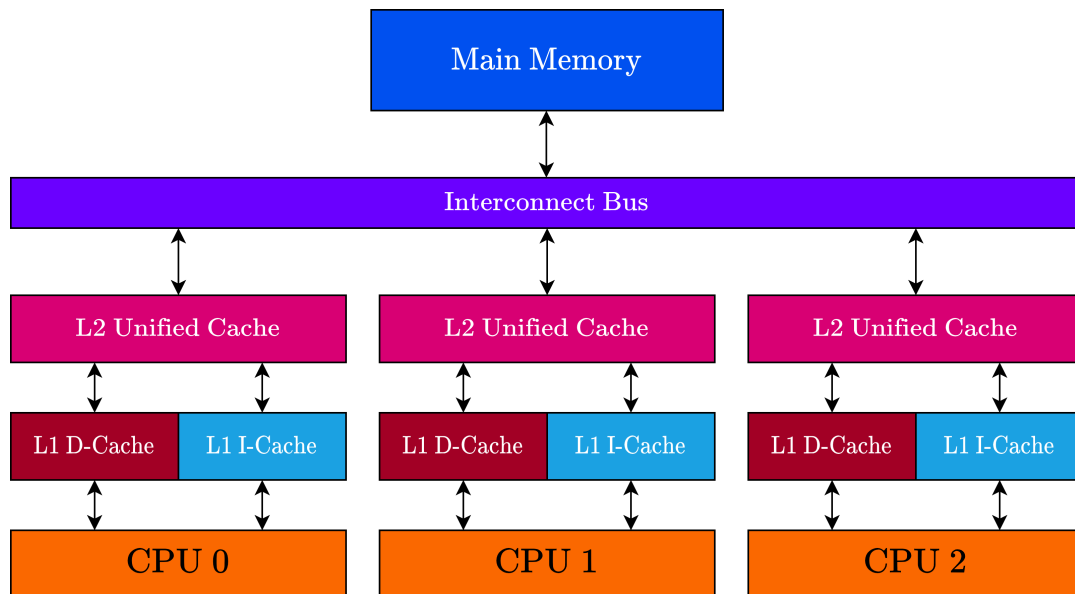
```
        END DO
    END DO
    DO j=1, m
        DO i=1, n
            A(i,j) = B(i,j)
        END DO
    END DO

    ! Complete communication
    DO i=1, 4
        CALL MPI_WAIT(req(i), status(1.i), ierr)
    END DO
...
END DO
```

8.4   Snooping Cache Coherency Protocols

Given a value can be in multiple caches, at multiple levels

- Which is the latest version of the cache line?
- Which parts of the line are actually outdated?
- Can the cached copy be used?

The goal is to ensure the result of any execution is the same as if operations in each thread were executed in a sequential order (sequential consistency), or something weaker but with enough guarantees to allow concurrent programs to be written correctly.

Sequential Consistency Never!	Extra Fun! 8.4.1
sequential consistency is a very strong memory model, and is not typically implemented by any architectures.	

8.4.1   Invalidation

Instead of updating values in all caches on a write, we instead invalidate.

- On the first write, send invalidation signal to other core's cache controllers.
- Hence after first write, we know we have the *only* copy, so no need to communicate further writes.
- Sharing state stored in extra bits added to the cache line.

This is typically faster than updating other cache entires on write, unless the data is usually required immediately.

A *snooping cache controller* is placed between the L2 cache and the interconnect bus to monitor for invalidations, and send invalidations.

8.4.2 Berkely Protocol

Each cache line contains a state:

Invalid	
Valid	Clean, potentially shared & unowned
Shared-dirty	modified, possibly shared, owned
Dirty	modified, not aliased/only copy, owned

On a read hit, a **clean** or **dirty** entry can be read (read from the owner), **invalid** requires an access on the interconnect bus and **shared-dirty** requires an invalidation to be sent.

Read Miss

```
Broadcast the request on the interconnect bus

if other cache has line as DIRTY or SHARED-DIRTY:
    get the cache line
    set its cache line to SHARED-DIRTY
    set our cache line to VALID
else:
    load line from main memory
    set our cache line to VALID
```
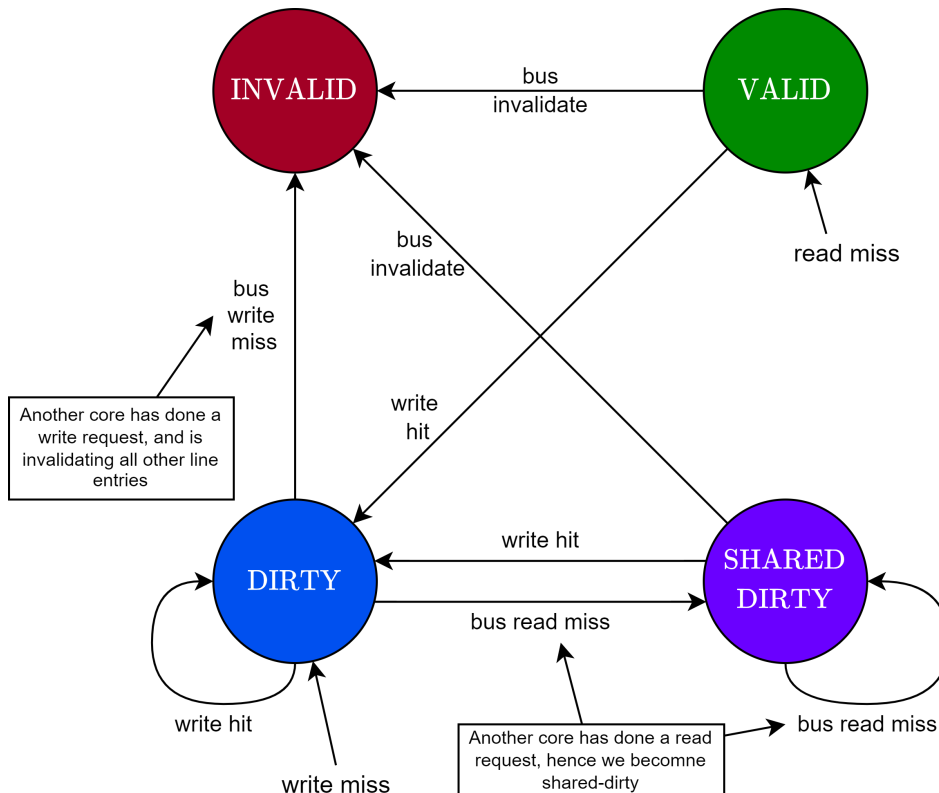
Write Hit

```
if cache line is VALID or SHARED-DIRTY:
    send invalidation to interconnect bus
    set our cache line to DIRTY
```

Write Miss

```
get line from owner
set all copies of the line to INVALID
set our cache line to DIRTY
```


There are alternatives ssuch as:

MESI Protocol	Clean for exclusive state (no miss for private data on write)
Illinois Protocol	Cache supplies data when shared state (no memory access)

The bus and CPU may contend for cache access:

- Can duplicate the tags of the L1 cache to allow CPU and snooping cache controller to check in parallel.

- Can use the L2 cache as a filer (L2 contains all of L2 - *multi-level inclusion*), hence bus can check L2 cache, if not present then a lien is also not present in L1.

- Creates constraints on the cache design.

8.5 Synchronisation

UNFINISHED!!!

8.6 Scalable Shared Memory

UNFINISHED!!!

- Can duplicate the tags of the L1 cache to allow CPU and snooping cache controller to check in parallel.

- Can use the L2 cache as a filer (L2 contains all of L2 - *multi-level inclusion*), hence bus can check L2 cache, if not present then a lien is also not present in L1.

Chapter 9

Asymptotic Approach

UNFINISHED!!!

Chapter 10

Credit

Image Credit

Front Cover   Intel i386 die shot by Pauli Rautakorpi on wikimedia here.

Content

Based on the architecture course taught by Prof Paul Kelly.

These notes were written by Oliver Killane.