

# 60009

Distributed Algorithms  
Imperial College London

# Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Course Structure & Logistics . . . . .	2
1.2	Course Resources . . . . .	2
1.3	Distributed Systems . . . . .	3
<b>2</b>	<b>Elixir</b>	<b>4</b>
2.1	learning Elixir . . . . .	4
2.2	The Elixir System . . . . .	5
2.3	Message Passing . . . . .	6
<b>3</b>	<b>Credit</b>	<b>8</b>

# Chapter 1

## Introduction

### 1.1 Course Structure & Logistics



**Dr Narankar Dulay**

The module is taught by Dr Narankar Dulay.

**Theory** For weeks 2  $\rightarrow$  10:

- Elixir (learning programming language)
- Introduction
- Reliable Broadcast
- FIFO, casual and total order Broadcast
- Consensus
- Flip Improbability Result
- Temporal Logic of Actions
- Modelling Broadcast
- Modelling Consensus

### 1.2 Course Resources

The course website contains all available slides and notes.

## 1.3 Distributed Systems

### Distributed System

### Definition 1.3.1

A set of processes connected by a network, communicating by message passing and with no shared physical clock.

- No total order on events by time (no shared clock)
- No shared memory.
- Network is logical - processes may be on the same OS process, same VM, same machine different machines communicating over a physical network.

Distributed systems must contend with the inherent uncertainty (failure, communication delay and an inconsistent view of the system's state) in communication between potentially physically independent processes (fallible machines, networks and software).

### Leslie Lamport

### *Extra Fun! 1.3.1*

A computer scientist and mathematician, credited with creating TLA (used on this course), as well as being the initial developer of latex (used for these notes).

” There has been considerable debate over the years about what constitutes a distributed system. It would appear that the following definition has been adopted at SRC:

A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable. ”

# UNFINISHED!!!

# Chapter 2

## Elixir

### 2.1 learning Elixir

- [Introduction To Elixir & Installation](#)
- [Elixir Documentation and Standard Library](#)
- [Elixir Learning Resources](#)
- [Devhints Exlixir Cheatsheet](#)
- [Elixir Quick Reference](#)
- [Learn Elixir in Y Minutes](#)

#### Two Sum

#### Example Question 2.1.1

Write a program to provide the two indexes of numbers in a list that sum to a given target. (This is the famous leetcode problem two sum).

```
defmodule Solution do
  @spec two_sum(nums :: [integer], target :: integer) :: [integer]
  def two_sum(nums, target) do
    nums
    |> Enum.with_index()
    |> Enum.reduce_while(%{}, fn {num, idx}, acc ->
      case Map.get(acc, target - num) do
        nil ->
          { :cont, Map.put(acc, num, idx) }
        val ->
          { :halt, [idx, val] }
      end
    end)
  end
end
```

We could also write this recursively with a helper function

```
defmodule Solution do
  @spec two_sum(nums :: [integer], target :: integer) :: [integer]
  def two_sum(nums, target) do
    two_sum_aux(nums, target, %{}, 0)
  end

  defp two_sum_aux([next | rest], target, prevs, index) do
    val = Map.get(prevs, target - next)
    if val != nil do
      [val, index]
    else
      two_sum_aux(rest, target, Map.put(prevs, next, index), index + 1)
    end
  end
end
```

```

    end
  end
end

```

## Add two numbers

## Example Question 2.1.2

Given The following linked list structure, write a program taking two numbers (represented in reverse as linked lists), and produce a linked list of their sum. (This is leetcode problem add two numbers)

*# Definition for singly-linked list.*

```

defmodule ListNode do
  @type t :: %__MODULE__{
    val: integer,
    next: ListNode.t() | nil
  }
  defstruct val: 0, next: nil
end

```

```

defmodule Solution do
  @spec add_two_numbers(l1 :: ListNode.t | nil, l2 :: ListNode.t | nil) :: ListNode.t | nil
  def add_two_numbers(l1, l2) do
    x = get_list(l1) + get_list(l2)
    if x == 0 do
      %ListNode{val: 0, next: nil}
    else
      to_list(x)
    end
  end

  defp get_list(node) do
    case node do
      %ListNode{val: v, next: n} -> v + 10 * get_list(n)
      nil -> 0
    end
  end

  defp to_list(n) do
    case n do
      0 -> nil
      i -> %ListNode{val: rem(i,10), next: to_list(div(i,10))}
    end
  end
end

```

## 2.2 The Elixir System

### Elixir

### Definition 2.2.1

A concurrent (with actors) and functional programming language used for fault tolerant distributed systems.

- A modernized successor language to Erlang
- Runs using BEAM (Erlang's virtual machine) and hence compatible with erlang
- Has many additions over erlang (protocols, streams and metaprogramming)

A lightweight user level thread (green threads) managed by the runtime.

- Everything is a process.
- Processes are strongly isolated, when two processes interact it does not matter which nodes, or even machines they run on.
- Processes share no resources (cannot share variables), they can only interact through message passing.
- Process creation and destruction is fast.
- Processes interact by message passing.
- Processes have unique names, if a name is known it can be used to pass messages
- Error handling is non-local.
- Processes do what they are supposed to do or fail.

All elixir processes run within a node, a node can manage many processes (creation, scheduling, and garbage collection).

- A node runs as an OS process, potentially with several OS threads scheduled across several cores.
- Multiple nodes can run on a single machine (or virtual machine such as a docker container).
- A node can efficiently manage thousands to millions of elixir processes.

Communication between processes is implemented through shared memory on the same machine and TCP when over a network. However processes are not exposed to this - the same primitives are used for inter and intra node/machine communication.

## 2.3 Message Passing

The `send` and `receive` statements are used for message passing.

```
# send somedata (any type) to process p
send p, somedata

# Wait until a message that matches the pattern is added to the message queue
# (or a timeout occurs), then remove it (potentially skipping over messages
# that do not match)
receive do
  somepattern -> dosomething(somepattern)
  # ... some other patterns
end
```

- Each process has its own message queue.
- Messages received are appended to the message queue of the receiving process.
- The sender does not wait for the message to be appended, it continues immediately after sending.

We can implement a basic client-server system in this way. Here we are using a component-based approach (split the program into components, each asynchronously message pass), by convention each component is an elixir module, modules can be instantiated in many processes & (by convention) have a public `start()` function.

```
defmodule Cluster do
  def start do
    # Spawn two processes, with the function start
    # Server.ex and Client.ex are modules containing a public start function
    # (Assuming we have started a client_node and server_node)
    s = Node.spawn(:'server_node@172.19.0.2', Server, :start, [])
    c = Node.spawn(:'client_node@172.19.0.1', Client, :start, [])
  end
end
```

```

# We send the PIDs of the processes to each other, we can pattern match on
# atoms for convenience in receiving
send s, { :bind, c }
send c, { :bind, s }
end
end

```

```

defmodule Server do
  def start do
    receive do
      { :bind, c } -> next(c)
    end
  end

  # next is defined as private, here
  # recursion is used for iteration.
  # To avoid a stack overflow tail
  # recursion is required
  defp next(c) do
    receive do
      { :circle, radius } ->
        send c, { :result, 3.14 * radius
                  * radius }

      { :square, side } ->
        send c, { :result, side * side }
    end
    next(c)
  end
end

```

```

defmodule Client do
  def start do
    receive do
      { :bind, s } -> next(s)
    end
  end

  defp next(s) do
    send s, { :circle, 1.0 }
    receive do
      { :result, area } ->
        IO.puts "Area is #{area}"
    end
    Process.sleep(1000)
    next(s)
  end
end

```



# Chapter 3

# Credit

## Image Credit

## Content

Based on the distributed algorithms course taught by Prof Narankar Dulay.

These notes were written by Oliver Killane.