

60017

System Performance Engineering
Imperial College London

Contents

1	Introduction	2
1.1	Logistics	2
1.1.1	Extra Resources	2
1.2	What is System Performance Engineering	2
1.3	Performance Engineering Process	3
1.3.1	Metrics	3
1.3.2	Quality of Service (QoS) Objectives	4
1.3.3	Service Level Agreements	4
1.4	Performance Evaluation Techniques	4
1.4.1	Measuring	4
1.5	Optimisation Loop	5
2	Profiling	7
2.0.1	Call Stack Tracing	8
2.1	Sampling	8
2.1.1	Time-Base Intervals	9
2.1.2	Event-Base Intervals	9
2.1.3	Indirect Tracing	9
2.2	Recording Events	10
2.2.1	Manual Instrumentation	10
2.2.2	Automatic Instrumentation	10
2.2.3	Binary Instrumentation	10
2.2.4	Kernel Counters/ Software Performance Counters	11
2.2.5	Emulation	11
2.2.6	Hardware Counters	11
2.3	Perf	11
2.4	Microarchitectural bottleneck analysis	11
2.5	Vtune	12
3	Modelling	13
3.1	Motivation	13
3.2	Numerical Models	13
3.3	Analytical Models	14
3.3.1	Empirical to Analytical	14
3.4	Memory Access Patterns	15
3.4.1	Memory Region	15
3.4.2	Sequential	16
3.4.3	Repetitive Random Access	16
3.4.4	Random Traversal	17
3.4.5	Modelling Complex Patterns	17
3.5	Modelling State	17
3.5.1	Simple Branch Direction Prediction	18
3.6	Modelling an Entire System	20
4	Efficient Code	21
4.1	Motivation	21
4.2	CPU Efficiency	21
5	Credit	22

Chapter 1

Introduction

1.1 Logistics



Dr Holger Pirk



Dr Luis Vilanova

First Half

- Hardware efficiency in complex systems
- Scaling up
- "getting the most bang for buck"

Second Half

- Scale out Processing

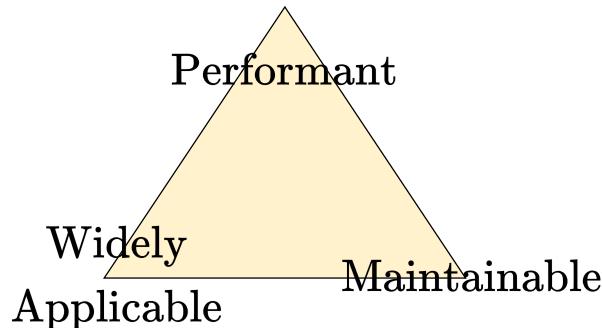
1.1.1 Extra Resources

UNFINISHED!!!

1.2 What is System Performance Engineering

System	Definition 1.2.1
A collection of components interacting to achieve a greater goal. <ul style="list-style-type: none">• Usually applicable to many domains (e.g a database, operating system, webserver). The goal is domain-agnostic• Designed to be flexible at runtime (deal with other interacting systems, real conditions) (e.g OS with user input, database with varying query volume and type)• Operating conditions are unknown at development time (Database does not know schema prior, OS does not know number of users prior, Tensorflow does not know matrix dimensionality prior) <p>Large & complex systems are typically developed over years by multiple teams.</p>	

The challenge with *system performance engineering* is to make systems maintainable, widely applicable and fast.



System Performance Engineering

Definition 1.2.2

Performance engineering encompasses the techniques applied during a systems development life cycle to ensure the non-functional requirements for performance will be met.

- Functional requirements (correctness, features) are assumed to be met.
-

High Performance Computing

Definition 1.2.3

High performance programming uses highly distributed & parallel computer systems (e.g supercomputers, clusters) to solve advanced problems.

- Focuses on solving a single computationally difficult problem.
- Workloads are well defined and known at development time.
- Sometimes supported by custom hardware (e.g FPGAs, ASICs, custom CPU extensions)

1.3 Performance Engineering Process

1.3.1 Metrics

A *target metric* is used to quantitatively measure any improvement in *performance* (e.g for use in a *SLA*). The metric needs to be well defined:

- When measuring starts (e.g when to measure latency from)
- Where measuring is done (is server response time measured on server, on a client, under what conditions?)

Imperials

Example Question 1.3.1

Provide some examples of metrics regarding a database.

Latency	Measuring time to query, planning time, the whole system's response time over a network.
Throughput	Measure the maximum requests/second possible (often used to compare web-servers)
Memory Usage scalability	measurable, but must be careful (e.g. OS interaction) Can define a metric regarding how quickly some metric (e.g. throughput) increases with scale (e.g. instances of a distributed system)

It is also important to define when a requirement is satisfied.

- Setting an optimisation budget (e.g. in developer hours)
- Setting a target or threshold (e.g. $x\%$ over baseline implementation)
- Combination of both

1.3.2 Quality of Service (QoS) Objectives

Quality of Service Objectives	Definition 1.3.1
A set of statistical properties of a metric that must hold for a system. <ul style="list-style-type: none">• Can include preconditions (e.g to define the environment/setup)• Can be in conflict with functional requirements (e.g framerate vs realism in graphics)	
Game On	Example Question 1.3.2
Give an example of a basic QoS Objective for a game's framerate. The game's framerate will be on average (over ... <i>preconditions</i> ...) 60fps if run on a GPU rated at 50GFlops or higher.	

1.3.3 Service Level Agreements

Service Level Agreements (SLAs)	Definition 1.3.2
Legal contracts specifying <i>QoS objectives</i> and penalties for violation. <ul style="list-style-type: none">• Non-functional requirements (not about system correctness)• Can be legally enforced	
Amazon	Extra Fun! 1.3.1
Amazon Web Services (AWS) provides a set of <i>service level agreements</i> relating to performance and availability. Violations are resolved by providing customers with service credits. Amazon SLAs	

When defining requirements for an SLA:

- | | |
|-------------------|---|
| Specific | State exact acceptance criteria (numerical terms). |
| Measurable | Ensure the metrics used can actually be measured. |
| Acceptable | Requirements should be rigorous such that meeting them is a meaningful success. |
| Realisable | Counter to Acceptable - need to be lenient enough to allow implementation. |
| Thorough | All necessary aspects of the system are specified. |

1.4 Performance Evaluation Techniques

1.4.1 Measuring

- Performed on the actual system (can be prototype or production/final).
- Can be difficult and costly (need to mitigate any impact of the measuring system on the system itself).
- As it is on the actual system, it can (if done properly) yield accurate results.

The two main types of measurement are:

Monitoring	Definition 1.4.1	Benchmarking	Definition 1.4.2
Measuring in production to get real usage performance metrics. <ul style="list-style-type: none">• Observe the system in its production environment• Collect usage statistics and analyse data (e.g user's preferred query types/structure, schema designs for databases)• Can monitor for and report SLA violations.		Measuring system performance in a controlled setting (e.g lab). <ul style="list-style-type: none">• The system to set into a predefined (or steady/hot) state• Perform some workload while measuring performance metrics.	

Benchmarking requires representative workloads in order to get metrics likely to be representative of a production environment.

Batch Workload

Definition 1.4.3

Program has access to entire batch at start of the benchmark.

- Useful when a throughput metric is being measured
- Simple to generate, and can even be recorded from a production environment.

Interactive Workload

Definition 1.4.4

A program generates requests to pass to the system being benchmarked.

- Useful when a latency metric is being measured.
- Workload generator needs to fast enough to saturate system being benchmarked.
- Often more representative of a production environment (e.g an operating system receives a workload over time)

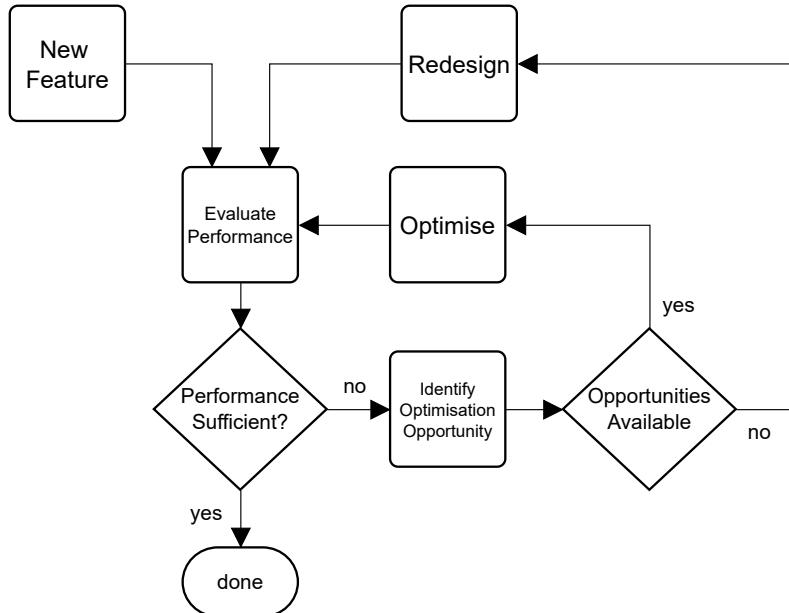
Hybrid

Definition 1.4.5

A common setup combining batch and interactive workload strategies (e.g sample random queries from a predefined work set).

In order to get useful results from which

1.5 Optimisation Loop



Performance Parameters

Definition 1.5.1

System and workload characteristics that affect performance.

System Parameters Do not change as the system runs (instruction costs, caches)

Workload Parameters Change as the system runs (available memory, users)

Numeric Parameters Quantitative (e.g CPU frequency, available memory, number of user)

Nominal Parameters Qualitative parameters (Runs on battery, has a GPU, runs in a VM)

The term *resource Parameters/Resources* refers to the parameters of the underlying platform (e.g CPU, memory).

Utilisation	Definition 1.5.2	Bottleneck	Definition 1.5.3
<p>The proportion of a resource used by a service to perform a service by a system.</p> <ul style="list-style-type: none"> • A service has limited resources available (e.g. CPU time, memory capacity, network bandwidth etc) • Total available resources/resource budget available to a service is a parameter 		<p>The resource with the highest utilisation.</p> <ul style="list-style-type: none"> • The limiting factor in performance of a system • given some resource x is the bottleneck, the system is x-bound (e.g. CPU-bound). • Not always a resource, and performance may be bottlenecked by some other factor (e.g. latency-bound → the system is dominated by waiting for some operation) 	

It is typically infeasible to identify all bottlenecks for an entire complex software system.

To limit optimisation complexity efforts should be restricted to optimising code paths that have particularly large effect on performance.

Critical Path	Definition 1.5.4	Hot Path	Definition 1.5.5
<p>The sequence of activities which contribute the largest overall duration.</p>		<p>A code path where most of the execution time is spent (e.g. very commonly executed subroutine)</p>	

In order to optimise we require:

- Ability to quickly compare alternative designs
- Ability to select a near optimal value for platform parameters

While workload parameters are not typically controllable at this stage, some system parameters are.

Parameter Tuning	Definition 1.5.6
<p>Finding the vector within the parameter space that minimises resource usage, or maximises performance.</p> <ul style="list-style-type: none"> • Exploring the parameter space is expensive (even with non-linear optimisation) • Analytical models can be used to accelerate search. • Tuning needs to consider tradeoffs (e.g. much of a cheap resource versus little of an expensive one) 	

Analytical Performance Model	Definition 1.5.7
<p>A model describing the relationship between system parameters and performance metrics.</p> <ul style="list-style-type: none"> • Having an accurate analytical model for the system is analogous to <i>understanding</i> the performance of the system. • Models can be stateless (e.g. an equation) or stateful (e.g. using markov chains). • Need to model dynamic systems with (ideally small) static models. • Very fast (faster than search parameter space) • Allow for <i>what-if analysis</i> of system and workload parameters. 	

Simulation	Definition 1.5.8
<p>A single observed run of a stateful model</p> <ul style="list-style-type: none"> • Can see all interactions within the system in perfect detail • Extremely expensive to run (limiting number of simulations and the speed of the optimisation workflow) • Much more rarely used than other techniques described 	

Chapter 2

Profiling

Event	Definition 2.0.1
A change in the state of the system.	
	<ul style="list-style-type: none">• Usually some granularity limit is used (e.g clock tick)• Optionally has a payload (properties describing the event - e.g cache line evicted → the addresses & data evicted)• Has an accuracy - degree to which the event represents reality (many events are numeric & come with measurement related error)
Simple/Atomic Event	Executed instruction, clock tick, function called
Complex Event	Cache line evicted, ROB flush due to misspeculation
Event sources have two components:	
Generator	Observes changes to system state (online → part of the runtime system)
Consumer	Processes events and converts into meaningful insights (offline or online)
The overhead associated with collecting events can be very high.	

Perturbation	Definition 2.0.2
The effect of analysis on the performance of a system.	
<ul style="list-style-type: none">• If it is constant/deterministic we can subtract it from measurement to get an accurate result.• Non-deterministic perturbation negatively affects accuracy as it cannot separate analysis overhead from measured performance	

Stacking up!	Example Question 2.0.1
Instrumentation is added to a program to inspect its stack at regular intervals, and record the stack trace. Assuming it is implemented to ensure the time spent traversing the trace is deterministic and can be removed from any results, why may this instrumentation still result in non-deterministic perturbation?	
Data Cache Side Effects	
The instrumentation may bring <i>deep</i> parts of the stack (<i>shallow</i> is <i>hot</i> and likely cached) into cache, evicting other lines. Hence the tracing <i>pollutes</i> the cache and results in more misses & hence more memory related stalls for the program.	
We could also more weakly argue about instruction cache limitations, or potential for associativity conflicts occurring may have been avoided by design in the original program, but due to placement of instrumentation's static data & text have their positions moved.	

Fidelity

Definition 2.0.3

The level of granularity with which events are recorded.

- Perfect fidelity means every event is recorded.
- Lower fidelity generally means less overhead.

Trace

Definition 2.0.4

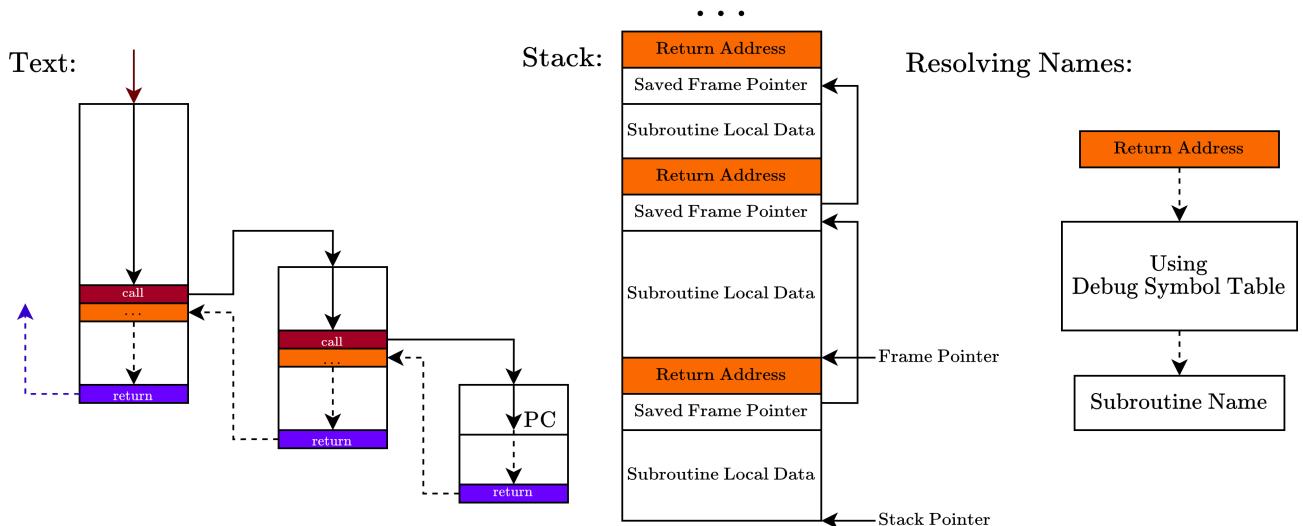
A complete log of the system's state (and hence changes in state/events) for some time period.

- Events usually totally ordered (exceptions being with parallelism - i.e multicore systems)
- Accuracy is inherited from events (e.g accuracy of clock, hardware counters etc)
- Analysing traces is time consuming (solution: Profiling)

Events may be aggregated to reduce the logging overhead.

2.0.1 Call Stack Tracing

A common trace is to capture and inspect the stack of a thread.



- Stack frame layout must conform to a convention that stores frame pointers and return addresses on the stack.
- Debug symbols included in the binary (e.g part of the ELF spec) can be used to convert return addresses scraped from the stack into symbols from the source code (function names).
- The stack may not easily represent the call structure of a program (inlining, tail call elimination, constant evaluation and propagation)

2.1 Sampling

Rather than recording all events, we can reduce overhead by sampling some events.

Performance	Fidelity traded against perturbation for reduced overhead in recording events.
Less Perturbation	Fewer interactions with program → less effect on performance → less perturbation.

Imperfect May skip sampling some events (e.g very small functions), but more expensive functions are more likely to be sampled, and we generally care about these more.

2.1.1 Time-Based Intervals

We can make use of hardware supported timer interrupts to set timers before using an interrupt to trap and allow the sampler to take control.

- Can use CPU *reference cycles* as a proxy metric to time
- Hardware clocks are often poorly defined & vary (i.e variable clock frequencies in modern CPUs, synchronisation of clocks between CPUs)
- Easily available & easy to interpret (ticks \propto time)

2.1.2 Event-Based Intervals

A generalisation of time-based intervals (as a clock tick \rightarrow time is an event)

- Can define in terms of occurrences of an event (e.g function call)
- Depending on design, can be accurate / low noise
- Can be tricky to interpret as a link to some measure proportional to time is needed (typically we care about execution time)

Quantisation Errors

Definition 2.1.1

The resolution of an interval is limited (e.g by clock), but time is continuous. Mapping events to a discrete measure of time can introduce errors & bias (costs may be attributed to the wrong time & hence wrong state).

2.1.3 Indirect Tracing

Indirect Tracing

Definition 2.1.2

We can sample from parts of a program, and infer traces from the structure of the program.

For example control flow instructions dominate non-control flow instructions, we could sample from control flow instructions, and then infer event between (non-control flow instructions) (called *basic block counting*), effectively indirectly tracing them.

- Can be used to reduce overhead from tracing
- fidelity and accuracy good (depending on the events traced and the indirection used)

Profile

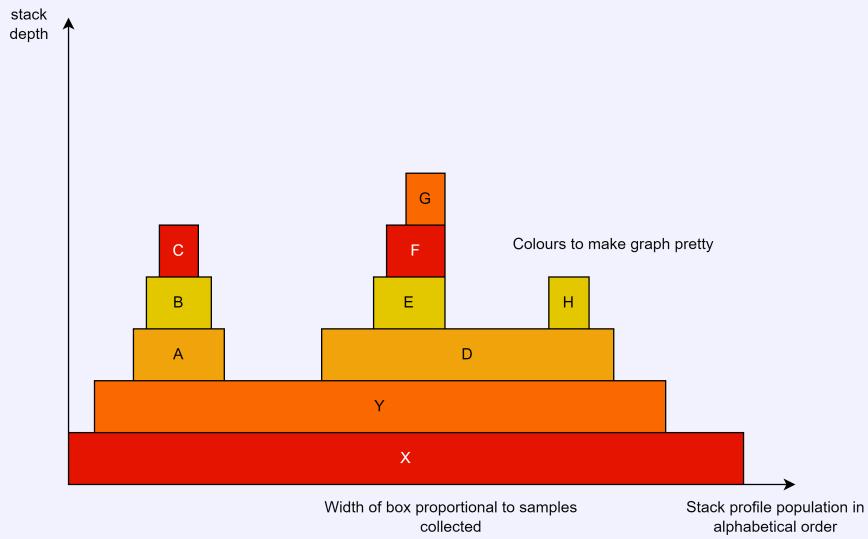
Definition 2.1.3

A (usually graphical) representation of information relating to the characteristics of a system in terms of resources (quantified) used in certain states.

- An aggregate over a specific metric (e.g a global aggregate if total cache misses, or per event such as cycles per instruction, or cache misses per function)
- Information is lost in aggregation
- Aggregation has lower overhead than recording all events, so can reduce perturbation

Flame Graph

Definition 2.1.4



2.2 Recording Events

Instrumentation

Definition 2.2.1

Adding event logging code to a program.

Easily Applicable No need to extra hardware or OS support.
Flexible Can implement any kind of logging required.

High Overhead
Perturbation As part of the program can effect performance

2.2.1 Manual Instrumentation

Logging using a library (e.g `printf` logging)

Fine control Programmer can easily specify exactly where to log what.
Supportless No need for compiler or hardware support.

High Overhead
Disable Need to disable for release builds (recompile without any logging)

2.2.2 Automatic Instrumentation

Compiler supported injection of event recording code into a program.

- Can be done at source level, or within some intermediate form (e.g injecting into some bytecode)
- Can potentially reduce overhead compared with manual instrumentation (compiler instruments at a lower level representation)

2.2.3 Binary Instrumentation

Instrumenting an already compiled binary.

Static	Adding instrumentation directly, overhead can be assessed from the binary.
Dynamic	Adding instrumentation at runtime (works well with JiT)

2.2.4 Kernel Counters/ Software Performance Counters

The kernel already has the tools required to collect many kinds of events. These tend to be higher-level OS interactions, rather than microarchitectural events. For example:

- Network packets sent
- Virtual memory events (e.g page faults)
- Context Switches
- Threads spawned

2.2.5 Emulation

2.2.6 Hardware Counters

Special registers configured to count low-level events as well as intervals (e.g cycles)

- Fixed number can be active at runtime
- Often buggy / unmaintained / inaccurate (and poorly documented) → only the most popular counters are trustworthy

2.3 Perf

UNFINISHED!!!

RTFM

Extra Fun! 2.3.1

Intel 64 and IA-32 Architectures Optimization Reference Manual.

- Microarchitectural features documented.
- Written as an optimisation guide.
- Code example can be found in its github repo

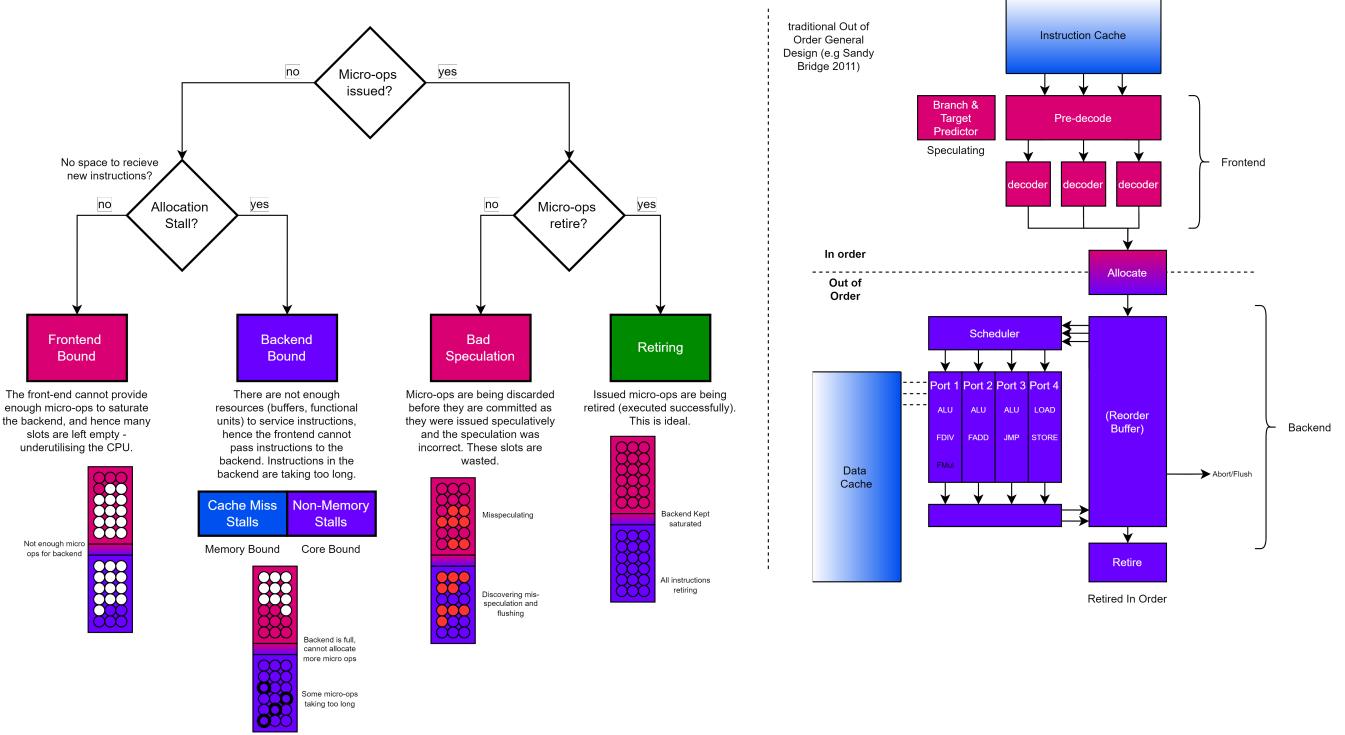
2.4 Microarchitectural bottleneck analysis

Advanced Computer Architecture

Extra Fun! 2.4.1

A basic understanding of computer architecture is required (pipelining, in order, caches, out of order, speculation).

The 60001 - Advanced Computer Architecture module by Prof Paul Kelly covers this in great depth.



See Intel's V-Tune performance metrics definitions for more

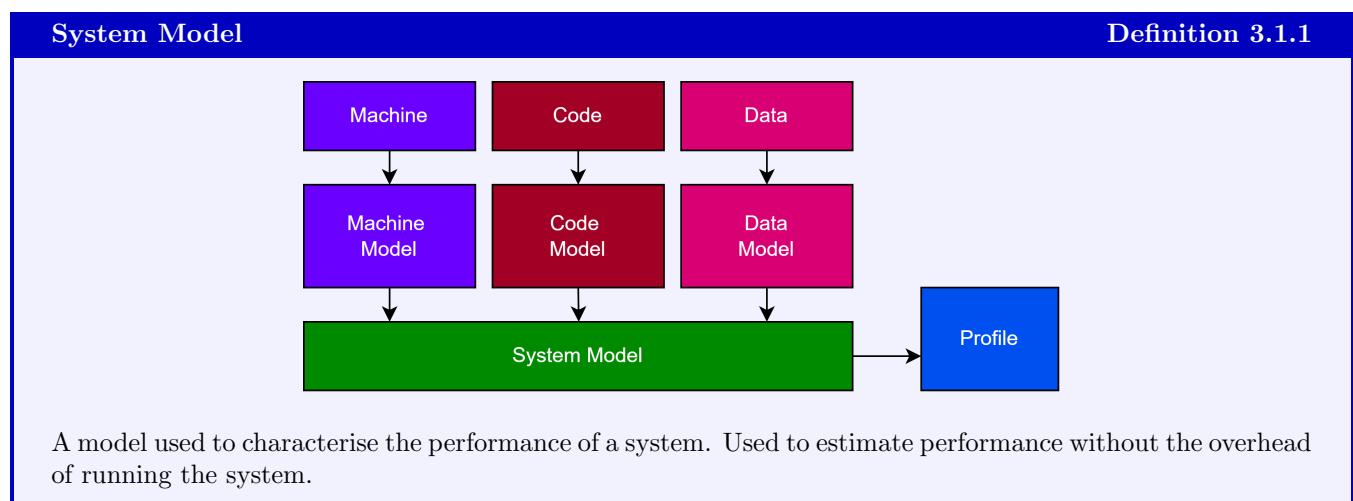
2.5 Vtune

UNFINISHED!!!

Chapter 3

Modelling

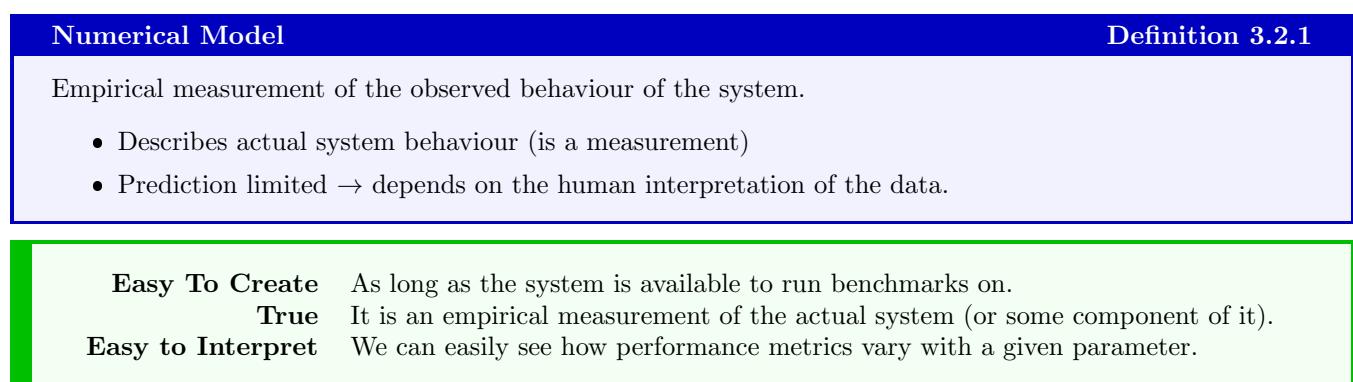
3.1 Motivation



During this course we make several simplifying assumptions:

- **Input Data from a Known Distribution**
Typically uniform. We avoid complications from correlated inputs.
- **No System Noise** Noise caused by the operating system (scheduling, other processes actions) and other factors.
- **Single Threaded & Deterministic Code** Modelling parallel systems requires considering contention and is an open area of research.

3.2 Numerical Models



Poor Generalisation	We cannot easily apply measurements to new values for parameters (mainly descriptive model / poor prediction)
Costly	For a large number of parameters, high fidelity a large amount of data is required, and hence many runs of benchmarks.
Limited Prediction	We can infer predictions from measurement, but it is difficult to extrapolate with confidence.

Microbenchmark

Definition 3.2.2

Small programs designed to test a specific portion of a system.

Numerical Models are constructed by:

1. Data (performance metrics) gathered through *microbenchmarks* on a range of parameter values.
2. Data is analysed/interpreted.

3.3 Analytical Models

Analytical Model

Definition 3.3.1

A formalised relationship between system parameters and performance metrics.

- Hard to interpret (limited descriptive use)
- Evaluating the model (given parameters) predicts performance metrics
- A detailed understanding of the system is required
- Model must be validated using experimental data

Models can even be used inside the system itself. For example using an analytical model to determine which optimiser to use for a query plan

Example Models

Extra Fun! 3.3.1

An example of a cost model for an R-Tree can be found in Cost models for join queries in spatial databases.

The Picasso Database Query Optimizer Visualizer another example of analytical models being used in the context of databases.

3.3.1 Empirical to Analytical

We fit an analytical model via regression on data from a numerical model.

Benchmark

For example we could benchmark the memory system of a machine:

```
#include <stdint.h>
#include <stddef.h>

extern int32_t* input_data; // An array of data
extern size_t N;           // A large constant
extern size_t size;        // Parameter: size of the region accessed

void benchmark() {
    volatile int32_t sum = 0;
    for (size_t i = 0; i < N; i++)
        // mod is used (expensive), to reduce use size power of 2 and a bitmask
        sum += input_data[i % size];
}
```

System Parameters

B_0	Size of a General Purpose Register of the CPU
B_1	Size of a cache line of the Level 1 cache
B_2	Size of a cache line of the Level 2 cache
B_3	Size of a Memory Page
l_0	Access Latency of the Level 1 Cache
l_1	Access Latency of the Level 2 Cache
l_2	Access Latency of the main memory
l_3	Lookup time in the Page Table
C_0	Capacity of a General Purpose Register of the CPU
C_1	Capacity of the Level 1 Cache
C_2	Capacity of the Level 2 Cache
C_3	Number of Memory Pages in the TLB \times Page size

Model

Given s is the stride parameter in the characteristic equations:

$$T_{Mem} = \sum_{i=0}^3 l_i \times \min\left(1, \frac{s}{B_i}\right) \text{ or we could model as } T_{Mem} = \begin{cases} l_0 & s < C_1 \\ l_0 + l_1 & s < C_2 \\ l_0 + l_1 + l_2 & s < C_3 \\ l_0 + l_1 + l_2 + l_3 & \text{otherwise} \end{cases}$$

We can compare the effects of altering system parameters on both models.

Fitting

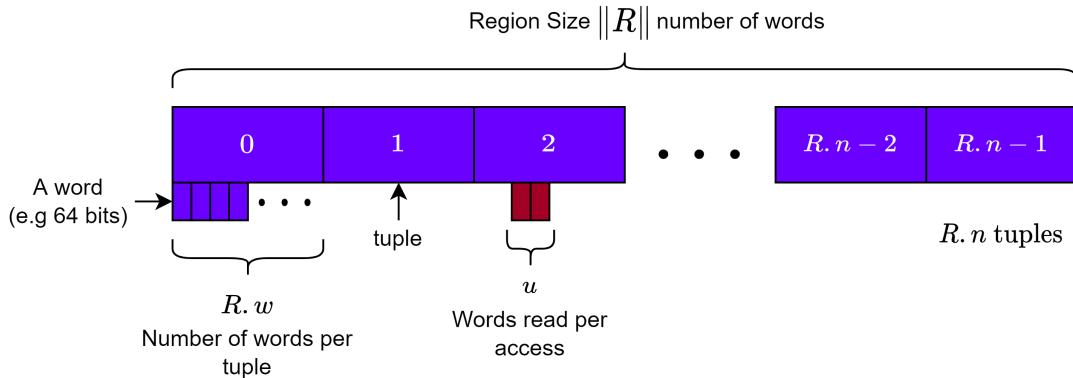
Some parameters can be taken from documentation, others may be derived from experiment.

Ideally the model should automatically self-tune to set parameters. Automating the tuning process has several advantages:

- Less work required by humans.
- Can adapt to changing conditions (e.g if cache size is artificially reduced by contention for shared cache by many threads)
- Can scale forward (if CPU is updated to newer generation, model can be re-tuned to fit changed system parameters)

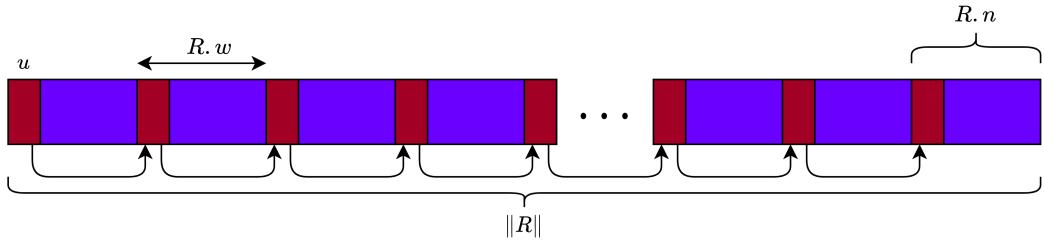
3.4 Memory Access Patterns

3.4.1 Memory Region



- The number of words skipped over between each access is $R.w - u$

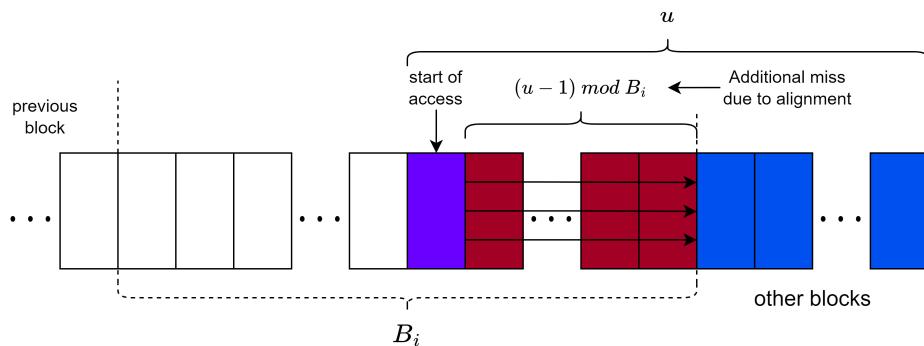
3.4.2 Sequential



Given some block size we can estimate the number of misses with a simple model.

- Assume a cold start (no part of the region in cache already)
- Assume accesses are within blocks, not over boundaries (e.g. if $u > 1$ and went over the edge of a cache line)
- s_trav means *sequential traversal*

$$M_i^s(s_trav) = \begin{cases} \frac{\|R\|}{B_i} & R.w - u < B_i \\ R.n \times \left\lceil \frac{u}{B_i} \right\rceil & R.w - u \geq B_i \end{cases}$$

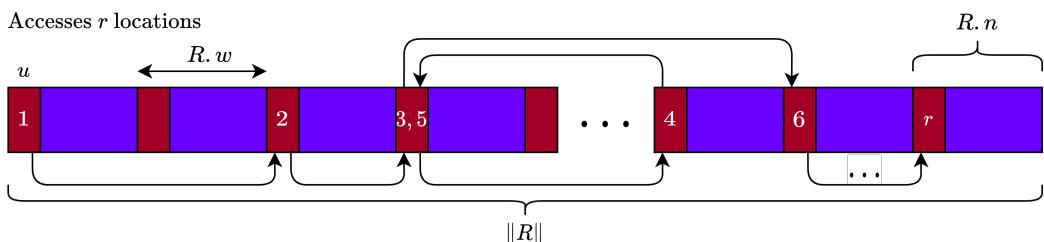


We may want to consider additional misses due to misalignment (u stretching over two B_i), in which case we must change the $R.w - u \geq B_i$ case to:

$$M_i^s(s_trav) = R.n \times \left(\left\lceil \frac{u}{B_i} \right\rceil + \frac{(u - 1) \bmod B_i}{B_i} \right)$$

3.4.3 Repetitive Random Access

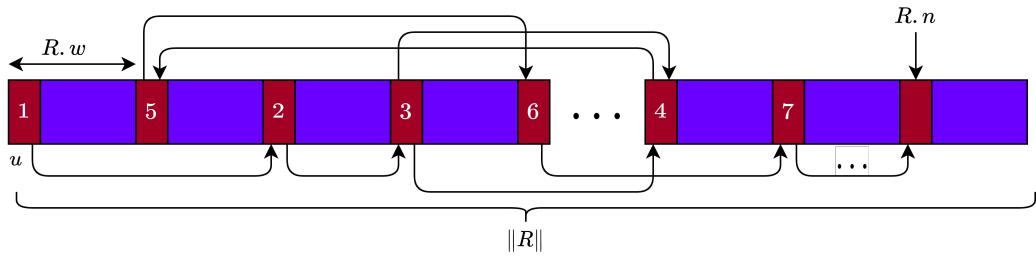
$$M_i^s(rr_acc) = \text{undefined in lecture}$$



Randomly access r locations (can repeat accesses)

3.4.4 Random Traversal

$$M_i^s(r_trav) = \text{undefined in lecture}$$



Access all $R.n$ locations in some random order.

3.4.5 Modelling Complex Patterns

- $\mathcal{P}_1 \oplus \mathcal{P}_2$ Sequential execution of patterns \mathcal{P}_1 then \mathcal{P}_2 .
- $\mathcal{P}_1 \odot \mathcal{P}_2$ Concurrent execution (access patterns interleaved).

We can hence combine access patterns.

Random Arrays	Example Question 3.4.1
Create an access pattern description for the following program. List any other assumptions.	
<pre>#include <stdint.h> #include <stddef.h> extern int32_t input_data_1[i]; // uniform random data (used for index) extern int32_t input_data_2[j]; // random data extern size_t N; // A large constant void benchmark() { int32_t sum = 0; for (size_t i = 0; i < N; i++) sum += input_data_2[input_data_1[i]]; }</pre>	
We assume that $N = i$ and all entries of <code>input_data_1</code> x are such that $0 \leq x \leq i$ (otherwise array accesses can be out of bounds)	
$s_trav(R.w = 1, u = 1, R.n = i) \odot rr_acc(R.w = 1, u = 1, R.n = j, r = i)$	

More complex access patterns can be modelled.

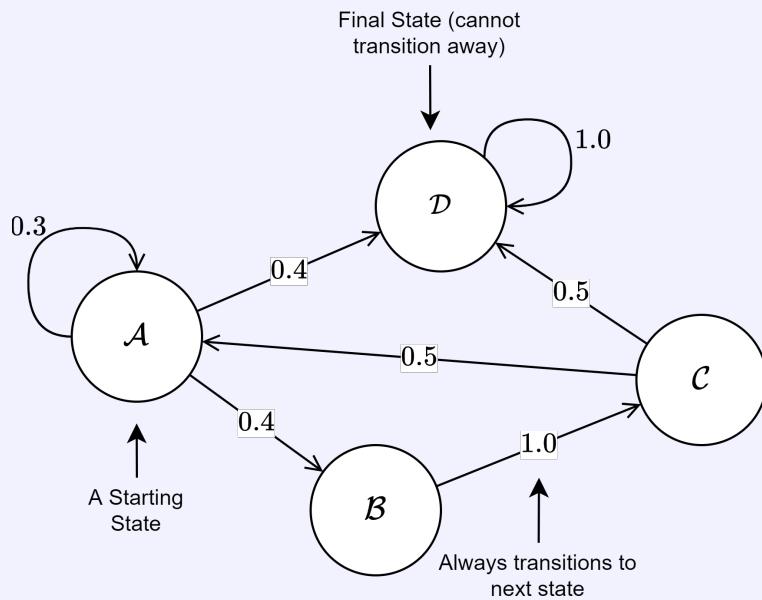
$$\underbrace{s_trav(\dots) \odot r_trav(\dots)}_{\text{Hash Join Build}} \oplus \left(\underbrace{\begin{array}{ccc} s_trav(\dots) \odot rr_acc(\dots) & \odot & s_trav(\dots) \\ \text{Read Input} & \text{Access Hashtable} & \text{Write Output} \end{array}}_{\text{Hash Join Build}} \right)$$

3.5 Modelling State

Analytical models are stateless, but most systems have some kind of dynamic state that can influence behaviour and performance.

Stochastic models can be used to encode state (combine analytical models and dynamic state).

A stochastic model describing a sequence (discrete time steps) of possible states.



- Much like a probabilistic finite state machine → the next state is only dependent on the previous (and random variables for transition) (called the *markov property*).
- Transition probabilities out of a state must sum to 1
- Can be represented as a matrix (directed graph represented as an adjacency matrix)

3.5.1 Simple Branch Direction Prediction

Consider a simple conditional branch, to a known target (instruction). For example a basic loop:

```
int N = 42;

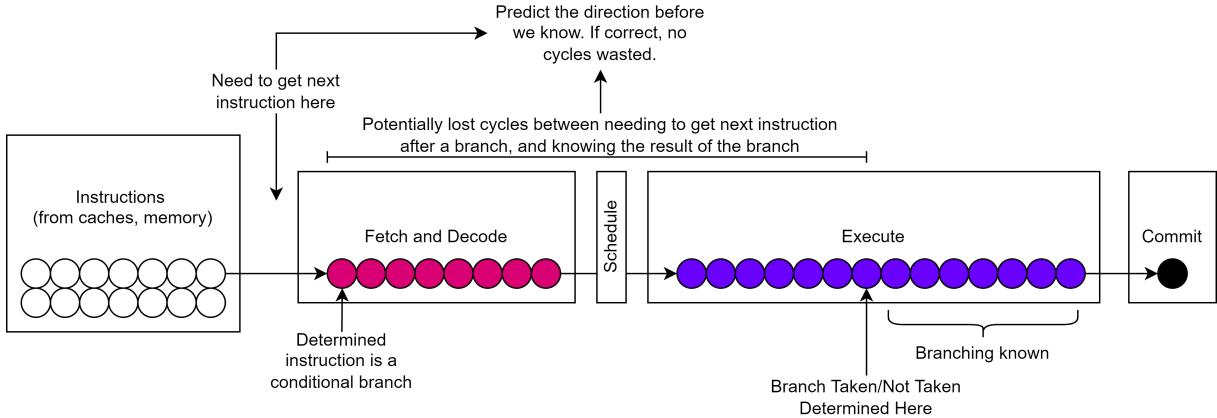
void a() {
    int i = 0;
    while (i < N) {
        i++;
    }
}
```

Compiled for x86 with -O1 (godbolt) we get the following:

```
a:
; Load N to edx and determine if any loop iteration should run (N may be 0)
mov    edx, DWORD PTR N[rip]
test   edx, edx

; Start looping
jle   .L1           ; [Conditional branch, known target]
mov    eax, 0         ; int i = 0
.L3:
add    eax, 1         ; i++
cmp    eax, edx       ; Check that i < N by checking if i == N yet
jne   .L3           ; [Conditional branch, known target]
.L1:
ret                ; [Conditional branch, unknown target -> target prediction with RAS]
N:
.long   42
```

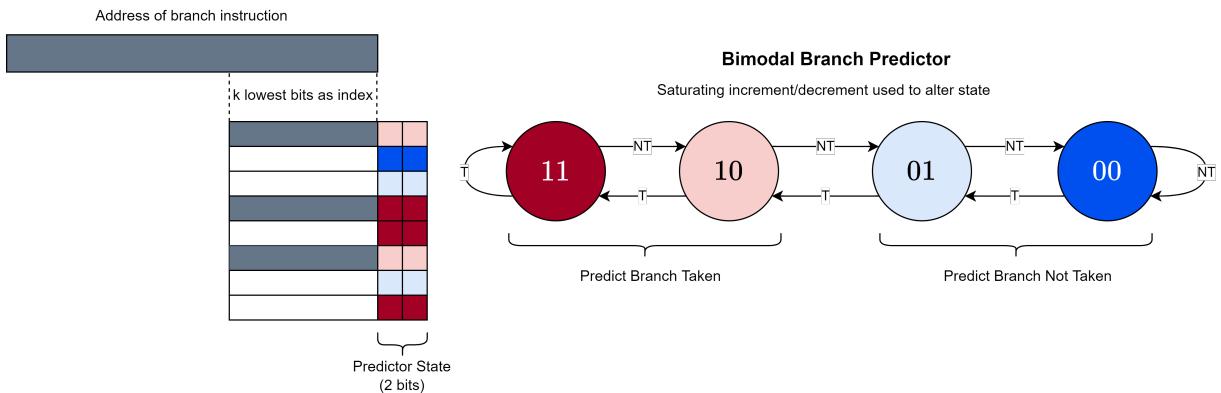
We will focus on the conditional branches with a known target (target prediction is more complex).



There are several ways to reduce potential penalty from branch instructions:

- Smaller frontend (reduce cycles between fetch and branch determination).
- Early branch resolution.
- Branch delay slot (branch does not take effect until n instructions after) (used in MIPS).
- Branch prediction.

A simple branch prediction scheme is to use a 2-bit branch history table, the result of branch resolution for a given branch instruction changes the state in the table, this state is used to predict other execution of branches at that instruction. We can represent the branch predictor state for a given instruction by a finite state machine.



Hence we can also model the branch predictor state as a *markov chain*.

Stationary Distribution

Definition 3.5.2

A probability distribution for a specific *markov chain* that remains the same (stationary) as time progresses.

Given the transition matrix \mathbf{P} of a chain, the stationary distribution can be represented by a vector of probabilities for being in any given state π_∞ such that:

$$\mathbf{P}\pi_\infty = \pi_\infty$$

Hence π_∞ is an eigenvector of \mathbf{P} with associated eigenvalue 1.

- It describes the probability of being in each state after infinite steps.

		Actual Takeness	
		Taken	Not Taken
Prediction	Taken	Correct	Mispredict
	Not Taken	Mispredict	Correct

$$P(\text{Mispredict}) = P(\text{Predict Taken}) \times P(\text{Actually Not Taken}) + P(\text{Predict Not Taken}) \times P(\text{Actually Taken})$$

3.6 Modelling an Entire System

Modelling an entire system is typically infeasible (scale and noise). Instead specific components and paths can be modelled.

1. Identify code that matters for performance using a profiler (Hot code sections, called bottlenecks in vtune)
2. Re-create behaviour in a controlled environment (*microbenchmarking*)
3. Create an analytical model of the code path/component.
4. Validate the model with experimental results.

Chapter 4

Efficient Code

4.1 Motivation

4.2 CPU Efficiency

CPU Bound	Definition 4.2.1
UNFINISHED!!!	
Control Hazard	Definition 4.2.2
UNFINISHED!!!	UNFINISHED!!!
Structural Hazard	Definition 4.2.3
UNFINISHED!!!	
Data Hazard	Definition 4.2.4
UNFINISHED!!!	

Chapter 5

Credit

Content

Based on the System Performance Engineering course taught by Dr Holger Pirk and Dr Luis Vilanova.

These notes were written by Oliver Killane.